Energy-Guided Smoothness for Robust Graph Classification under Label Noise

Anonymous Author(s)
Affiliation
Address
email

Abstract

Graph Neural Networks (GNNs) are vulnerable to label noise in graph classification, yet, unlike for image or node classification, the mechanisms that drive this vulnerability are poorly understood. We show that GNN robustness in this setting is governed by a property of the *learned representations*, not of the labels themselves: the within-graph Dirichlet energy of node features. Energy decreases while the model captures genuine class structure, but increases sharply, particularly in high-frequency components, when the model begins to memorize corrupted graph-level labels. We give a theoretical justification for why this decrease is non-trivial in graph classification by combining (i) a spectral argument that relates dataset-level energy to per-graph low-frequency structure and (ii) an entropy-based argument adapted from Qian et al. [2025], showing that uncontrolled energy growth disproportionately harms graphs that are most informative for classification. Building on this insight, we propose three complementary interventions, a spectral weight constraint (CE+W2), a Dirichlet-energy regularizer ($\mathcal{L}_{DE}$), and a robust loss function (GCOD), which all suppress harmful high-frequency energy through different mechanisms. Across eight benchmarks, including the full and modified ogbg-ppa, under symmetric noise (up to 60%), asymmetric noise, and a transfer to node classification on Cora, the three methods yield consistent gains while preserving clean-data performance. GCOD matches or exceeds the strongest noise-robust baselines (SOP, OMG) at $\sim$70$\times$ lower training overhead, establishing Dirichlet energy as both a diagnostic signal and a principled design target for noise-robust GNNs.

1 Introduction

Graph Neural Networks (GNNs) are now the workhorse of *graph classification*, the task of assigning a single label to an entire graph, with applications spanning molecular property prediction [Wale and Karypis, 2006], bioinformatics [Borgwardt et al., 2005], and social-network analysis [Yanardag and Vishwanathan, 2015a]. In real deployments, the labels attached to graphs are routinely noisy: assays disagree, annotators err, and weak supervision injects systematic mistakes. Although learning under label noise has been thoroughly studied for images [Algan and Ulusoy, 2021] and for node classification [Dai et al., 2021, Yuan et al., 2023a], graph classification under noisy labels is dramatically less explored [NT et al., 2019, Yin et al., 2023].

Graph classification is a fundamentally different setting. In node classification, label noise corrupts *individual node labels* on a single graph, and homophily, the tendency of connected nodes to share labels, is a structural prior that directly couples Dirichlet energy of the *label signal* to noise. In graph classification, label noise flips the label of an *entire graph*: nodes have no individual

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

classification labels, and within-graph homophily/heterophily is largely orthogonal to the graph-level label. The standard tools and intuitions from node classification therefore do not transfer.

Our central observation is non-trivial. We measure the Dirichlet energy of *learned node representations*, not of labels. Empirically, this energy first decreases as the model captures genuine, transferable patterns. Then, when the model begins to memorize incorrect graph-level labels, the within-graph energy rises sharply, with the surge concentrated in the high-frequency spectrum. This is not a tautology: noise is injected only at the graph-label level, while node features and edges are unchanged. To fit incorrect labels *after* pooling, the network must restructure its node-level representations, and the only way to do so without losing the correctly labelled signal is to amplify high-frequency directions. The connection between graph-level label noise and within-graph spectral sharpening therefore emerges through training dynamics, not by definition.

Why does suppressing this energy growth help? Combining a dataset-level spectral identity (Proposition 1) with an entropy-based argument adapted from Qian et al. [2025], we show in Section 6 that uncontrolled within-graph energy growth disproportionately damages *high-entropy graphs*, the very graphs that carry the most discriminative signal at the pooling stage. A non-trivial decrease of within-graph Dirichlet energy thus simultaneously (i) prevents the network from fitting graph-level label noise and (ii) preserves the high-entropy regions on which graph classification accuracy depends.

Contributions.

- We give the first systematic study of *when* GNNs fit graph-level label noise (Section 4), isolating three failure modes: over-parameterization relative to task complexity, low-order graphs, and small training sets.
- We establish Dirichlet energy as a diagnostic for graph-level noise memorization (Section 5), with a frequency decomposition showing the surge is concentrated in high-frequency components.
- We provide *theoretical grounding for why a non-trivial energy decrease prevents noise fitting in graph classification* (Section 6), through a dataset-level spectral identity (Proposition 1) and an entropy-based result (Theorem 1) showing that high-entropy graphs are the most sensitive to over-sharpening. This is what makes the connection *non-trivial* for graph classification, in contrast to the trivial relationship that exists for node classification.
- We translate the theory into three convergent interventions (Section 7): spectral weight constraints (CE+W2), explicit Dirichlet regularization ($\mathcal{L}_{DE}$), and a centroid-based loss GCOD. The fact that mechanistically distinct interventions all reduce energy and improve robustness provides causal evidence beyond a single-method study.
- We evaluate on eight benchmarks under symmetric, asymmetric and high-rate (60%) noise, the full and modified ogbg-ppa, and a node-classification transfer (Cora). GCOD matches or exceeds OMG [Yin et al., 2023] and SOP [Liu et al., 2022] at $\sim$70$\times$ lower compute, and our energy-based methods generalize beyond their original setting.

2 Related Work

Learning under label noise. Robust loss functions [Ghosh et al., 2017, Patrini et al., 2017], sample re-weighting [Liu and Tao, 2015, Liu et al., 2022], early-learning regularizers [Arpit et al., 2017], and small-loss filtering [Gui et al., 2021] dominate the image-classification literature. Wani et al. [2023] introduce NCOD, a centroid-based outlier-discounting loss that we adapt to graphs.

Noisy-label learning on graphs. Most prior work concerns *node* classification, exploiting homophily, contrastive losses, or pseudo-labels [Dai et al., 2021, Yuan et al., 2023a,b, Li et al., 2024, Kang et al., 2018]. For *graph* classification, NT et al. [2019] propose a surrogate loss that discards estimated-noisy labels under specific assumptions, and OMG [Yin et al., 2023] combines contrastive learning, MixUp [Lim et al., 2022], and curriculum filtering. We differ in two ways: (i) we provide a spectral/energy-based explanation of *why* GNNs fit graph-level noise, rather than only mitigating it, and (ii) our practical method (GCOD) achieves comparable robustness at a fraction of the compute.

Dirichlet energy and over-smoothing. Dirichlet energy quantifies signal smoothness on graphs [Zhou and Schölkopf, 2005]. Existing analyses establish that GNNs act as low-pass filters whose energy decays exponentially with depth under spectral conditions [NT and Maehara, 2019, Cai and Wang, 2020, Oono and Suzuki, 2020, Rusch et al., 2023], and that weight-matrix eigenvalues control attraction/repulsion of node features [Giovanni et al., 2023]. Energy-based regularizers [Zhou et al., 2021, Bo et al., 2021, Chen et al., 2023] address *too little* energy (over-smoothing) on *node* classification. We address the opposite regime: *too much* within-graph energy caused by graph-level label noise. Qian et al. [2025] recently show that, for graph classification specifically, over-smoothing in *high-entropy* regions is the most damaging; we leverage this insight (Theorem 1) to argue why suppressing energy growth is essential under graph-level noise. A more comprehensive related-work discussion is in Appendix A.

3 Background

Notation. A graph is $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathbf{X})$ with $N = |\mathcal{V}|$, degree matrix $\mathbf{D}$, adjacency $\mathbf{A} \in \{0,1\}^{N\times N}$, feature matrix $\mathbf{X} \in \mathbb{R}^{N\times m}$, and normalized Laplacian $\mathbf{\Delta} = \mathbf{I} - \mathbf{D}^{-1/2}\mathbf{A}\mathbf{D}^{-1/2}$. In *graph classification*, the dataset is $\mathcal{D} = \{(\mathcal{G}^i, \mathbf{y}_i)\}_{i=1}^n$, $\mathbf{y}_i \in \{0,1\}^{|C|}$ the one-hot graph label; we write c_i for the class. Under message passing, $\mathbf{H}_i^{l+1} = \mathrm{UP}_{\mathbf{\Omega}_l}(\mathbf{H}_i^l, \mathrm{AGGR}_{\mathbf{W}_l}(\mathbf{H}_i^l, \mathbf{A}^i))$, with $\mathbf{H}_i^0 \equiv \mathbf{X}^i$ and $\mathbf{Z}^i \equiv \mathbf{H}_i^L$. A permutation-invariant readout $f_\theta : \mathbb{R}^{N_i\times m'} \to \mathbb{R}^{|C|}$ maps node features to class probabilities $\hat{\mathbf{y}}_i$. Under noisy labels, the observed c_i may differ from the ground truth; the test set is clean.

Dirichlet energy on graphs. For graph $\mathcal{G}^i$ with representation $\mathbf{Z}^i \in \mathbb{R}^{N_i\times m'}$,

$$E^{\text{dir}}(\mathbf{Z}^i) = \sum_{(u,v)\in\mathcal{E}^i} \left\| \mathbf{Z}_u^i/\sqrt{d_u} - \mathbf{Z}_v^i/\sqrt{d_v} \right\|_2^2. \tag{1}$$

Small E^{dir} means smooth (low-frequency) representations within the graph; large E^{dir} means sharpening (high-frequency). To track behavior across a class-restricted subset $\mathcal{S}$ of training graphs we use

$$E_{\text{set}}^{\text{dir}}(\mathcal{S}) = \frac{1}{|\mathcal{S}|} \sum_{\mathcal{G}^i\in\mathcal{S}} E^{\text{dir}}(\mathbf{Z}^i). \tag{2}$$

Crucially, (1) is computed on *learned representations*, not on labels. Within-graph homophily of features therefore evolves with training, decoupled from the graph-level label of $\mathcal{G}^i$.

4 When GNNs Fit Graph-Label Noise

GNNs trained with standard cross-entropy (CE) on graph classification tasks are not uniformly vulnerable to label noise. On the full ogbg-ppa [Szklarczyk et al., 2018a] (37 classes, the standard OGB benchmark), even 40% symmetric noise barely degrades CE accuracy, because the model is under-parameterized for the full task and behaves robustly [Zhang et al., 2021]. To expose the noise-fitting regime, we control three axes: **(F1)** number of classes (proxy for over-parameterization), **(F2)** graph order (size of each graph), and **(F3)** training-set size. We use a 5-layer GIN [Xu et al., 2019] with 300 hidden units throughout, unless stated.

In all three regimes (Fig. 1 and Appendix D), GNNs progressively memorize noisy graph-level labels. This narrows the question to: what changes in the network's internal representations during memorization, and can this change be controlled?

5 Dirichlet Energy as a Diagnostic for Noise Fitting

We track $E_{\text{set}}^{\text{dir}}(\mathcal{D}_c)$, the within-graph energy averaged over training graphs of class c. Across architectures and datasets (Fig. 2), the diagnostic is consistent: $E_{\text{set}}^{\text{dir}}$ falls during the early "learning" phase, then rises precisely when the model begins to fit noisy graphs (CE-20% in Fig. 2b), with 40% noise yielding a strictly larger end-of-training energy than 10% noise (Appendix F).

Energy is sensitive to noise specifically through high-frequency components. A natural concern is that energy also rises during clean-data overfitting [Yan et al., 2022]; we confirm this in Appendix F but show the noise-driven surge is markedly steeper. To localize the surge spectrally, we use the HLFF-GNN decomposition [Xu et al., 2024] to split representations into low-frequency (Z_1) and

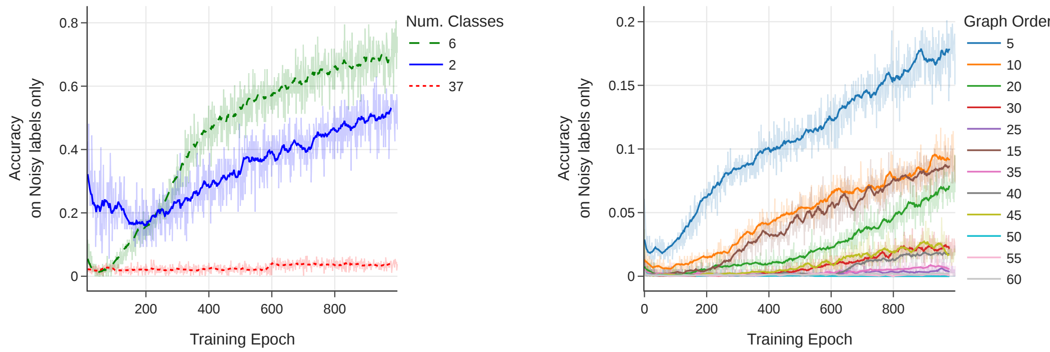

(a) Effect of class count.

(b) Effect of graph order.

Figure 1: Training accuracy on *noisy labels only*: a higher curve means more memorization. (a) Reducing the number of PPA classes (with 20% symmetric noise) makes a fixed GIN over-parameterized, accelerating noise fitting. (b) Smaller average graph order (synthetic data, 35% noise) leaves less internal structure for averaging, enabling memorization.

high-frequency (Z_2) components. On ENZYMES with 30% noise, $E^{\text{dir}}(Z_1)$ remains essentially unchanged, while $E^{\text{dir}}(Z_2)$ rises monotonically, and slope/variance statistics of $E^{\text{dir}}(Z_2)$ provide an early warning for label noise (Appendix E). HLFF-GNN is used purely as an analytical tool; the same total-energy surge appears in standard GIN/GCN.

Why energy is more useful than a generic overfitting signal. Validation loss tells you *that* the model is overfitting. Within-graph Dirichlet energy tells you *how*: which spectral components drive the failure. This mechanistic specificity is what turns a diagnostic into a design principle, and motivates the three interventions in Section 7, each of which targets the same harmful spectral behavior in a distinct way.

6 Why a Non-Trivial Energy Decrease Prevents Noise Fitting

We now justify why suppressing within-graph energy growth is a principled defense in graph classification. The argument has two parts: a dataset-level spectral identity that lets us reason about the whole training set as a single block-diagonal graph, and an entropy-based result that explains why uncontrolled energy growth disproportionately harms classification.

6.1 Dataset-Level Spectral Identity

Proposition 1 (Dataset-level Dirichlet energy)**.** *Let* $\mathcal{D} = \{\mathcal{G}^1, \ldots, \mathcal{G}^n\}$ *with representations* $\mathbf{Z}^i$ *and adjacencies* $\mathbf{A}^i$*. Define the block-diagonal graph* $\bar{\mathcal{G}} = (\bar{\mathbf{Z}}, \bar{\mathbf{A}})$ *with* $\bar{\mathbf{Z}} = [\mathbf{Z}^1 \| \ldots \| \mathbf{Z}^n]$ *and* $\bar{\mathbf{A}} = \operatorname{blkdiag}(\mathbf{A}^1, \ldots, \mathbf{A}^n)$*. Then*

$$E_{set}^{dir}(\mathcal{D}) = \frac{1}{|\mathcal{D}|} E^{dir}(\bar{\mathbf{Z}}) = \frac{1}{|\mathcal{D}|} \sum_{r=1}^{m'} \sum_{u=1}^{N_{tot}} \bar{\lambda}_u \, (\bar{\boldsymbol{\psi}}_u^\top \bar{\mathbf{Z}}_{:,r})^2, \tag{3}$$

where $\{(\bar{\lambda}_u, \bar{\boldsymbol{\psi}}_u)\}$ *are eigenpairs of* $\bar{\boldsymbol{\Delta}}$*, and* $\bar{\Lambda} = \bigcup_i \Lambda^i$ *is the union of per-graph spectra.*

The proof, given in Appendix B, exploits $\det(\bar{\boldsymbol{\Delta}} - \lambda \mathbf{I}) = \prod_i \det(\boldsymbol{\Delta}^i - \lambda \mathbf{I})$. The implication is conceptually important: *decreasing* $E_{set}^{dir}(\mathcal{D})$ *during training is equivalent to enhancing the low-frequency components of the corresponding block-diagonal graph.* Smoothing therefore acts *simultaneously* on all training graphs, in proportion to their per-graph spectra.

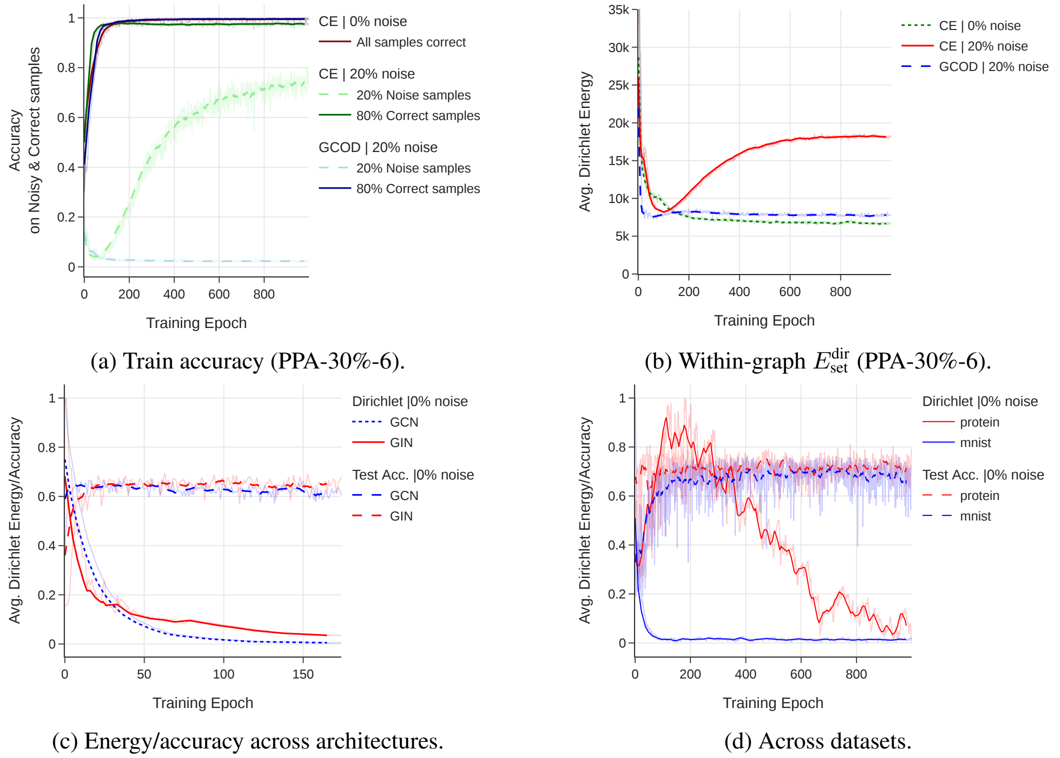

(a) Train accuracy (PPA-30%-6).

(b) Within-graph $E^{\text{dir}}_{\text{set}}$ (PPA-30%-6).

(c) Energy/accuracy across architectures.

(d) Across datasets.

Figure 2: Within-graph Dirichlet energy of learned representations *rises during noise memorization*, across architectures (GIN/GCN) and datasets (PPA, MUTAG, MNIST, PROTEINS). Under clean labels, $E^{\text{dir}}_{\text{set}}$ steadily decays as accuracy improves; under noise, it follows the same decay until the model starts fitting noisy graphs, after which it sharply rises. GCOD suppresses this surge.

6.2 An Entropy-Based Argument for Graph Classification

We now combine Proposition 1 with the structural-entropy framework of Qian et al. [2025]. For a graph $\mathcal{G}$, the structural Shannon entropy is

$$H(\mathcal{G}) = -\sum_{i=1}^{|V|} \frac{d_i}{\sum_j d_j} \log \frac{d_i}{\sum_j d_j}. \tag{4}$$

Qian et al. [2025] prove (i) that loosely-connected graphs with more nodes have higher H (their Theorem 1), and (ii) that, at fixed $|V|$, denser connectivity yields strictly higher H (their Theorem 2). High-entropy graphs are precisely those whose pooled representation depends on a richer mixture of frequencies, since their normalized Laplacian spectrum spreads over a wider range [Qian et al., 2025]. Combining this with our spectral identity yields:

Theorem 1 (Non-trivial role of energy decrease in graph classification)**.** *Let $\bar{\mathcal{G}}$ be the block-diagonal graph from Proposition 1, and let $\mathcal{G}^i, \mathcal{G}^j$ be two non-isomorphic graphs in $\mathcal{D}$ with $H(\mathcal{G}^j) \geq H(\mathcal{G}^i)$. Suppose the GNN sharpens representations along high-frequency directions of $\bar{\mathbf{\Delta}}$, i.e. E^{dir}_{set} grows. Then the increase $\Delta E^{dir}(\mathbf{Z}^j) \geq \Delta E^{dir}(\mathbf{Z}^i)$, and the discriminability of the pooled representations (measured by JSD on the softmax of $\mathbf{Z}^j, \mathbf{Z}^i$) collapses faster for the higher-entropy graph.*

Proof sketch (full proof in Appendix B). By Proposition 1, sharpening lifts spectral mass into eigenvectors $\bar{\psi}_u$ with large $\bar{\lambda}_u$. Each such eigenvector concentrates on the connected component (i.e. a single graph) whose Laplacian contributes that eigenvalue. By Theorems 1–2 of Qian et al. [2025], the high-frequency tail of Λ^i is denser for higher-$H(\mathcal{G}^i)$, so the high-entropy graph absorbs more of the spectral mass and its ΔE^{dir} dominates. The JSD argument follows from the entropic over-smoothing analysis of Qian et al. [2025, Sec. 3.3], which links over-sharpening of high-entropy regions to a collapse of node-distribution distinctness, and therefore of the pooled graph descriptor. $\square$

What this gives us, and why it is non-trivial. For *node* classification, the link between Dirichlet energy and noise is essentially trivial: noise corrupts the label signal directly, the label signal is defined on the same nodes, and homophily ties Dirichlet energy of the labels to noise by definition. *None of these shortcuts apply in graph classification*: the corrupted signal is a single graph-level label, the energy is computed on learned features of nodes that have no individual labels, and within-graph homophily is unrelated to the graph's class. Theorem 1 closes this gap by showing that the network can only fit incorrect graph-level labels by sharpening high-frequency, high-entropy regions, the same regions whose preservation is most important for graph classification accuracy. A non-trivial decrease of $E^{\text{dir}}_{\text{set}}$ therefore prevents fitting to noise *and* preserves classification capacity, an alignment that does not exist in the node-classification regime.

Remark 1. This result also explains an empirical observation in Section 5: the surge of $E^{\text{dir}}(Z_2)$ under noise is much steeper than under clean overfitting. Clean overfitting can be absorbed across the spectrum, while noise fitting is forced into the high-frequency, high-entropy directions identified above.

7 Three Convergent Methods

The theory of Section 6 identifies three distinct levers to suppress harmful within-graph energy growth: the *spectrum of the weight matrices* (architectural), the *energy of representations* (regularization), and the *loss function* (sample re-weighting). We design one intervention for each, and use them as convergent evidence for the energy framework.

7.1 Method 1: Spectral Constraint on Weight Matrices (CE+W2)

Giovanni et al. [2023] show that, in linear graph convolutions, positive eigenvalues of the post-aggregation weight matrix induce attraction (smoothing) between connected nodes, while negative eigenvalues induce repulsion (sharpening). Together with our Proposition 1, this suggests that removing negative eigenvalues from the post-aggregation matrix $\mathbf{W}_l^2$ should suppress harmful sharpening dataset-wide.

After each gradient step, we compute $\mathbf{W}_l^2 = \mathbf{\Phi}_l^2\, \boldsymbol{\mu}_l^2\, (\mathbf{\Phi}_l^2)^{-1}$, project $\boldsymbol{\mu}_l^2 \to [\boldsymbol{\mu}_l^2]^+ = \max(\boldsymbol{\mu}_l^2, 0)$, and reconstruct $\mathbf{W}_l^{2+} = \mathbf{\Phi}_l^2 [\boldsymbol{\mu}_l^2]^+ (\mathbf{\Phi}_l^2)^{-1}$. The projection is treated as a deterministic post-update operation, not part of the gradient graph. **CE+W2** is CE training with this projection; it requires no clean validation data and no extra hyperparameters. It is most useful as a theoretical instantiation of the energy story; in practice, the eigendecomposition is costly and can destabilize convergence (Appendix C.4), motivating the alternatives below.

7.2 Method 2: Explicit Dirichlet-Energy Regularization ($\mathcal{L}_{DE}$)

A more direct lever is to penalize $E^{\text{dir}}(\mathbf{Z}^i)$ above a threshold U_{c_i} that is class-aware:

$$\mathcal{L}_{DE}(\mathcal{D}) = \frac{1}{N} \sum_{i=1}^{N} \big[\max(0, E^{\text{dir}}(\mathbf{Z}^i) - U_{c_i})\big]^2, \quad \mathcal{L} = \mathcal{L}_{CE} + \lambda\, \mathcal{L}_{DE}. \tag{5}$$

We instantiate the threshold in two ways. **Class-specific:** U_c is recomputed every epoch as the mean E^{dir} over a small clean validation set restricted to class c. This adapts to class-dependent baseline energies (e.g. a small molecular graph's natural energy is different from a social graph's), but requires a clean held-out set. **Fixed:** $U_c = U$ globally, a single hyperparameter that does not require any clean validation data, at the cost of mild over-smoothing on clean labels. Both variants improve robustness; the class-specific variant additionally improves *clean-label* accuracy by acting as a generalization-aware regularizer.

7.3 Method 3: Graph Centroid Outlier Discounting (GCOD)

GCOD is our practical recommendation. It extends NCOD [Wani et al., 2023] (originally an image-classification loss) to graph classification, with two design changes that are necessary in the graph-classification regime: a third KL term that disentangles clean from noisy graphs via per-sample trainable parameters $\mathbf{u} \in [0, 1]^n$, and feedback through the current training accuracy a_{train} that prevents premature outlier discounting. For batch B with predictions $f_\theta(\mathbf{Z}_B)$, hard predictions $\hat{\mathbf{y}}_B$, soft labels $\tilde{\mathbf{y}}_B$ (centroid-based, see Appendix C.6), and observed labels $\mathbf{y}_B$:

$$\mathcal{L}_1 = \mathcal{L}_{CE}\big(f_\theta(\mathbf{Z}_B) + a_{\text{train}}\,\text{diag}_{\text{vec}}(\mathbf{u}_B)\,\mathbf{y}_B,\ \tilde{\mathbf{y}}_B\big), \tag{6}$$

$$\mathcal{L}_2 = \frac{1}{|C|}\big\|\hat{\mathbf{y}}_B + \text{diag}_{\text{vec}}(\mathbf{u}_B)\,\mathbf{y}_B - \mathbf{y}_B\big\|^2, \tag{7}$$

$$\mathcal{L}_3 = (1 - a_{\text{train}})\,D_{KL}\big\{\mathfrak{L},\ \sigma(-\log \mathbf{u}_B)\big\}, \tag{8}$$

with $\mathfrak{L} = \log \sigma(\text{diag}_{\text{mat}}(f_\theta(\mathbf{Z}_B)\mathbf{y}_B^\top))$. Updates are split: $\theta \leftarrow \theta - \alpha\nabla_\theta(\mathcal{L}_1 + \mathcal{L}_3)$, $\mathbf{u} \leftarrow \mathbf{u} - \beta\nabla_{\mathbf{u}}\mathcal{L}_2$. Empirically, GCOD reduces $E^{\text{dir}}_{\text{set}}$ during training (Fig. 2) without explicitly minimizing it: by down-weighting graphs flagged as noisy via large u_i, it removes the only gradient pressure that would push the network toward high-frequency directions in the first place. GCOD *is not* a theoretically derived energy minimizer. Instead, the theory predicts that any sample-level mechanism that suppresses noise-driven gradients will reduce $E^{\text{dir}}_{\text{set}}$, which is exactly the convergent evidence we observe.

8 Experiments

Setup. We evaluate on ogbg-ppa [Szklarczyk et al., 2018a] (full and a controlled 30%/6-class subset), MUTAG, PROTEINS, ENZYMES, IMDB-BINARY, REDDIT-MULTI-12K, MSRC-21, MNIST-graph, and Cora (node classification, transfer test). Backbones are GIN [Xu et al., 2019] and GCN [Kipf and Welling, 2017]; full hyperparameters are in Appendix C. Baselines: standard CE; SOP [Liu et al., 2022], a strong noise-robust loss from image classification; OMG [Yin et al., 2023], the leading graph-classification noise method. All numbers are mean$\pm$std over 4 runs.

Why a controlled PPA subset? On the full 37-class PPA, all methods are under-parameterized for the task, and CE itself is robust (Section 4); a noise-robustness method cannot help where the baseline does not fail. We use a 30%/6-class PPA subset as a *controlled* testbed where over-parameterization triggers noise fitting, and we additionally report the full-PPA result to show our methods do no harm there. Generality claims are grounded in the seven additional benchmarks.

8.1 Main Result: Controlled PPA Subset

Table 1 reports the full ablation on the controlled PPA subset. GCOD is best at every noise level and, crucially, has the smallest gap between best and final test accuracy, indicating that it does not overfit late in training. CE+W2 and $\mathcal{L}_{DE}$ both improve robustness over CE, confirming the convergent-evidence prediction of Section 6. Under clean labels, $\mathcal{L}^*_{DE}$ even outperforms CE, a direct effect of the class-aware energy threshold.

8.2 Generalization Across Datasets and Noise Types

Table 2 stress-tests GCOD on seven additional benchmarks, including the previously requested 60% symmetric noise regime. GCOD consistently improves over CE: $+5\%$ to $+8\%$ at 60% noise on most datasets, and matches CE under low or asymmetric noise. The robustness scales monotonically with noise rate (consistent with the energy-noise scaling in Appendix F).

Full ogbg-ppa. Under the standard 37-class benchmark, all methods perform within statistical noise of CE (CE: 66.4%, GCOD: 66.7% at 20% symmetric noise), confirming GCOD does no harm where baselines do not fail.

Comparison with OMG. Under the OMG [Yin et al., 2023] protocol (Table 3), GCOD matches OMG's accuracy (within 0.1–0.2 pp) while requiring $\sim$*70*$\times$ *less* training overhead (Table 4): OMG adds $\approx$200% runtime relative to CE-GIN, while GCOD adds only $\approx$2.9%. The accuracy-per-FLOP advantage is substantial and matters for any deployment beyond toy benchmarks.

Cross-setting transfer to node classification (Cora). As a stress test of generality, we apply GCOD and CE+W2 unchanged to node classification on Cora (20% symmetric noise), comparing against 13 node-specific methods (NRGNN, RTGNN, GNN-Cleaner, ERASE, CR-GNN, etc.). CE+W2 ranks $\#2$ (0.8157) on GAT and $\#5$ (0.7247) on GCN; GCOD ranks $\#4$ on both backbones (0.7910 GAT, 0.7403 GCN). Despite being designed for graph classification, both methods outperform several specialized node-classification approaches, evidence that the energy-based perspective captures something general about GNN training dynamics under label noise.

Table 1: Test accuracy on the PPA-30%-6 subset (5-layer GIN, mean±std over 4 runs). **Best** is highlighted in red and bold, second-best in blue. "CE clean only" shows accuracy on the clean subset of the noisy data, as a reference upper bound. Class-specific $\mathcal{L}_{DE}^*$ requires a clean validation set.

Noise	Method	Test Acc.		Train Acc.	
		Best	Last	Best	Last
0%	CE	96.25 ± 0.05	91.25	100.0	99.33
	GCOD	96.65 ± 0.52	93.25	99.23	99.03
	CE+W2	96.50 ± 1.12	86.50	99.76	99.29
	$\mathcal{L}_{DE}$ (Fixed)	96.58 ± 0.27	85.70	99.26	98.29
	$\mathcal{L}_{DE}^*$ (Class-spec.)	96.96 ± 0.04	88.67	99.41	98.67
20%	CE (clean only)	96.15 ± 0.09	90.66	84.50	84.07
	CE	88.66 ± 0.16	62.58	94.98	93.45
	SOP	91.01 ± 0.44	85.50	79.17	77.64
	GCOD	93.91 ± 0.26	92.58	80.11	79.52
	CE+W2	89.83 ± 1.23	76.66	84.90	83.68
	$\mathcal{L}_{DE}$ (Fixed)	89.69 ± 0.57	80.25	78.07	70.14
	$\mathcal{L}_{DE}^*$	92.34 ± 0.41	80.92	84.55	83.73
40%	CE (clean only)	95.08 ± 0.13	80.44	68.15	67.05
	CE	82.08 ± 0.22	51.83	82.47	78.88
	SOP	82.33 ± 1.32	65.41	58.90	57.68
	GCOD	93.88 ± 0.04	91.08	65.09	64.31
	CE+W2	88.25 ± 1.13	56.33	68.78	65.88
	$\mathcal{L}_{DE}$ (Fixed)	88.58 ± 0.75	77.12	61.08	54.53
	$\mathcal{L}_{DE}^*$	88.55 ± 0.11	71.83	68.98	67.80

Table 2: GCOD matches or beats CE across diverse datasets, including domains where the baseline is already strong. Mean±std over 4 runs. *Sym*$_{20}$*:* 20% symmetric. *Asym*$_{40}$*:* 40% asymmetric. *Sym*$_{60}$*:* 60% symmetric (high-noise stress test).

Dataset	0% CE	Sym$_{20}$ CE	Sym$_{20}$ GCOD	Asym$_{40}$ CE	Asym$_{40}$ GCOD	Sym$_{60}$ CE	Sym$_{60}$ GCOD	Δ_{60}
PROTEINS	81.16	76.23	79.38 ± 0.85	72.90	76.19 ± 1.22	61.42	69.87 ± 1.40	+8.45
MNIST	72.69	66.30	71.26 ± 0.44	71.15	72.61 ± 0.43	58.04	66.38 ± 0.61	+8.34
ENZYMES	73.33	64.16	68.69 ± 0.70	65.80	69.81 ± 0.98	50.21	58.12 ± 1.05	+7.91
IMDB-B	76.50	75.00	75.40 ± 0.92	68.18	72.89 ± 0.91	61.50	66.71 ± 1.08	+5.21
MUTAG	94.73	89.47	90.01 ± 0.38	91.16	93.19 ± 0.38	77.92	84.35 ± 1.55	+6.43
REDDIT	48.15	45.05	44.98 ± 2.87	48.01	47.94 ± 2.87	39.81	42.04 ± 2.45	+2.23
MSRC-21	96.69	90.26	94.69 ± 0.32	94.39	95.57 ± 0.32	79.46	87.92 ± 0.71	+8.46

8.3 Convergent Evidence for the Energy Story

The three methods are mechanistically distinct, CE+W2 is architectural, $\mathcal{L}_{DE}$ is regularization, GCOD is sample re-weighting, yet they all (i) reduce $E^{\text{dir}}_{\text{set}}(\mathcal{D}_c)$ during training (Fig. 2 and Appendix F) and (ii) improve robustness on noisy graph classification. This convergence is the strongest causal evidence we can offer in absence of a fully closed-form theoretical characterization: three independent levers acting on the same target produce the same outcome. Theorem 1 explains why this is expected.

9 Conclusion

We have argued, theoretically and empirically, that GNN robustness to graph-level label noise is governed by the within-graph Dirichlet energy of learned representations. The connection is

Table 3: OMG protocol (Yin et al., 2023). Percentage gain over CE.

Dataset	OMG	GCOD
MUTAG	0.061	0.062
IMDB-B	0.047	0.049
PROTEINS	0.039	0.041

Table 4: Relative training time (CE-GIN = 1.00).

Method	Runtime
CE-GIN	1.00
SOP	1.048
GCOD	**1.029**
CE+W2	1.33
OMG [Yin et al., 2023]	≈ 3.00

non-trivial in graph classification, requiring an entropy-based argument that does not exist in node classification, and it leads naturally to three convergent interventions (CE+W2, $\mathcal{L}_{DE}$, GCOD) that all improve robustness across eight benchmarks while preserving clean-label performance. GCOD, our practical recommendation, requires no clean validation data and adds only $\approx$3% training overhead, while matching the accuracy of state-of-the-art OMG at $\sim$70$\times$ lower compute.

Limitations. Our theoretical bridge to graph classification, while non-trivial, is not yet a closed-form characterization of the noise-energy relationship. Real-world noisy graph-classification benchmarks are scarce; we follow the standard synthetic-noise protocol, which limits ecological validity. Extending the framework to instance-dependent noise and to non-classification graph tasks (regression, link prediction) is a natural direction.

Broader impact. Robust graph learning matters in chemistry and biology, where annotation errors are systematic. Methods that achieve robustness without requiring clean validation sets (GCOD) reduce the cost of deploying GNNs in such settings.

References

Görkem Algan and Ilkay Ulusoy. Image classification with deep learning in the presence of noisy labels: A survey. *Knowledge-Based Systems*, 215:106771, 2021. ISSN 0950-7051. doi: https://doi.org/10.1016/j.knosys.2021.106771. URL https://www.sciencedirect.com/science/article/pii/S0950705121000344.

Eric Arazo, Diego Ortego, Paul Albert, Noel O'Connor, and Kevin Mcguinness. Unsupervised label noise modeling and loss correction. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 312–321. PMLR, 09–15 Jun 2019. URL https://proceedings.mlr.press/v97/arazo19a.html.

Devansh Arpit, Stanisław Jastrzębski, Nicolas Ballas, David Krueger, Emmanuel Bengio, Maxinder S. Kanwal, Tegan Maharaj, Asja Fischer, Aaron Courville, Yoshua Bengio, and Simon Lacoste-Julien. A closer look at memorization in deep networks. In Doina Precup and Yee Whye Teh, editors, *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 233–242. PMLR, 06–11 Aug 2017. URL https://proceedings.mlr.press/v70/arpit17a.html.

Deyu Bo, Xiao Wang, Chuan Shi, and Huawei Shen. Beyond low-frequency information in graph convolutional networks. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(5): 3950–3957, May 2021. doi: 10.1609/aaai.v35i5.16514. URL https://ojs.aaai.org/index.php/AAAI/article/view/16514.

Karsten M. Borgwardt, Cheng Soon Ong, Stefan Schönauer, S. V. N. Vishwanathan, Alex J. Smola, and Hans-Peter Kriegel. Protein function prediction via graph kernels. *Bioinformatics*, 21(suppl_1): i47–i56, 06 2005. ISSN 1367-4803. doi: 10.1093/bioinformatics/bti1007. URL https://doi.org/10.1093/bioinformatics/bti1007.

Chen Cai and Yusu Wang. A note on over-smoothing for graph neural networks, 2020. URL https://arxiv.org/abs/2006.13318.

Jialin Chen, Yuelin Wang, Cristian Bodnar, Rex Ying, Pietro Lio, and Yu Guang Wang. Dirichlet energy enhancement of graph neural networks by framelet augmentation, 2023. URL https://arxiv.org/abs/2311.05767.

Ching-Yao Chuang and Stefanie Jegelka. Tree mover's distance: Bridging graph metrics and stability of graph neural networks. *proceedings neurips*, 35:2944–2957, 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/file/139ae969f49abd9a113981c1f7fce5ce-Paper-Conference.pdf.

Enyan Dai, Charu Aggarwal, and Suhang Wang. Nrgnn: Learning a label noise resistant graph neural network on sparsely and noisily labeled graphs. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, KDD '21, page 227–236, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450383325. doi: 10.1145/3447548.3467364. URL https://doi.org/10.1145/3447548.3467364.

Enyan Dai, Wei Jin, Hui Liu, and Suhang Wang. Towards robust graph neural networks for noisy graphs with sparse labels. In *Proceedings of the Fifteenth ACM International Conference on Web Search and Data Mining*, WSDM '22, page 181–191, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450391320. doi: 10.1145/3488560.3498408. URL https://doi.org/10.1145/3488560.3498408.

Yair Davidson and Nadav Dym. On the holder stability of multiset and graph neural networks. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=P7KIGdgW8S.

Xuefeng Du, Tian Bian, Yu Rong, Bo Han, Tongliang Liu, Tingyang Xu, Wenbing Huang, Yixuan Li, and Junzhou Huang. Noise-robust graph learning by estimating and leveraging pairwise interactions. *Transactions on Machine Learning Research*, 2023. ISSN 2835-8856. URL https://openreview.net/forum?id=r7imkFEAQb.

James Fox and Sivasankaran Rajamanickam. How robust are graph neural networks to structural noise? *ArXiv*, abs/1912.10206, 2019. URL https://api.semanticscholar.org/CorpusID:209444207.

Fernando Gama, Joan Bruna, and Alejandro Ribeiro. Stability properties of graph neural networks. *IEEE Transactions on Signal Processing*, 68:5680–5695, 2019. URL https://api.semanticscholar.org/CorpusID:152282530.

Aritra Ghosh, Himanshu Kumar, and P Shanti Sastry. Robust loss functions under label noise for deep neural networks. In *Proceedings of the AAAI conference on artificial intelligence*, volume 31, 2017.

Francesco Di Giovanni, James Rowbottom, Benjamin Paul Chamberlain, Thomas Markovich, and Michael M. Bronstein. Understanding convolution on graphs via energies. *Transactions on Machine Learning Research*, 2023. ISSN 2835-8856. URL https://openreview.net/forum?id=v5ew3FPTgb.

Xian-Jin Gui, Wei Wang, and Zhang-Hao Tian. Towards understanding deep learning from noisy labels with small-loss criterion. In Zhi-Hua Zhou, editor, *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence, IJCAI-21*, pages 2469–2475. International Joint Conferences on Artificial Intelligence Organization, 8 2021. doi: 10.24963/ijcai.2021/340. URL https://doi.org/10.24963/ijcai.2021/340. Main Track.

Bo Han, Quanming Yao, Xingrui Yu, Gang Niu, Miao Xu, Weihua Hu, Ivor Tsang, and Masashi Sugiyama. Co-teaching: Robust training of deep neural networks with extremely noisy labels. *Advances in neural information processing systems*, 31, 2018.

Simona Juvina, Ana-Antonia Neacșu, Jean-Christophe Pesquet, Burileanu Corneliu, Jérôme Rony, and Ismail Ben Ayed. Training graph neural networks subject to a tight lipschitz constraint. *Transactions on Machine Learning Research Journal*, 2024.

Zhao Kang, Haiqi Pan, Steven C. H. Hoi, and Zenglin Xu. Robust graph learning from noisy data. *IEEE Transactions on Cybernetics*, 50:1833–1843, 2018. URL https://api.semanticscholar.org/CorpusID:56208557.

Jihye Kim, Aristide Baratin, Yan Zhang, and Simon Lacoste-Julien. Crosssplit: Mitigating label noise memorization through data splitting. In *International Conference on Machine Learning*, pages 16377–16392. PMLR, 2023.

Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. In *International Conference on Learning Representations*, 2017. URL `https://openreview.net/forum?id=SJU4ayYgl`.

Nils M. Kriege and Petra Mutzel. Subgraph matching kernels for attributed graphs. In *International Conference on Machine Learning*, 2012. URL `https://api.semanticscholar.org/CorpusID:8790929`.

Junnan Li, Richard Socher, and Steven C. H. Hoi. Dividemix: Learning with noisy labels as semi-supervised learning. In *ICLR*, 2020. URL `https://openreview.net/forum?id=HJgExaVtwr`.

Xianxian Li, Qiyu Li, De Li, Haodong Qian, and Jinyan Wang. Contrastive learning of graphs under label noise. *Neural Networks*, 172:106113, 2024. ISSN 0893-6080. doi: https://doi.org/10.1016/j.neunet.2024.106113. URL `https://www.sciencedirect.com/science/article/pii/S0893608024000273`.

Soon Hoe Lim, N. Benjamin Erichson, Francisco Utrera, Winnie Xu, and Michael W. Mahoney. Noisy feature mixup. In *International Conference on Learning Representations*, 2022. URL `https://openreview.net/forum?id=vJb4I2ANmy`.

Sheng Liu, Zhihui Zhu, Qing Qu, and Chong You. Robust training under label noise by over-parameterization. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu, and Sivan Sabato, editors, *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 14153–14172. PMLR, 17–23 Jul 2022. URL `https://proceedings.mlr.press/v162/liu22w.html`.

Tongliang Liu and Dacheng Tao. Classification with noisy labels by importance reweighting. *IEEE Transactions on pattern analysis and machine intelligence*, 38(3):447–461, 2015.

Marion Neumann, Roman Garnett, Christian Bauckhage, and Kristian Kersting. Propagation kernels: Efficient graph kernels from propagated information. *Machine Learning*, 102(2):209–245, 2016. ISSN 1573-0565. doi: 10.1007/s10994-015-5517-9. URL `https://doi.org/10.1007/s10994-015-5517-9`.

Hoang NT and Takanori Maehara. Revisiting graph neural networks: All we have is low-pass filters. *ArXiv*, abs/1905.09550, 2019. URL `https://api.semanticscholar.org/CorpusID:162183860`.

Hoang NT, Jun Jin Choong, and Tsuyoshi Murata. Learning graph neural networks with noisy labels, 2019. URL `https://openreview.net/forum?id=r1xOmNmxuN`.

Kenta Oono and Taiji Suzuki. Graph neural networks exponentially lose expressive power for node classification. In *International Conference on Learning Representations*, 2020. URL `https://openreview.net/forum?id=S1ldO2EFPr`.

Giorgio Patrini, Alessandro Rozza, Aditya Krishna Menon, Richard Nock, and Lizhen Qu. Making deep neural networks robust to label noise: A loss correction approach. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1944–1952, 2017.

Geoff Pleiss, Tianyi Zhang, Ethan Elenberg, and Kilian Q Weinberger. Identifying mislabeled data using the area under the margin ranking. *Advances in Neural Information Processing Systems*, 33: 17044–17056, 2020.

Feifei Qian, Lu Bai, Lixin Cui, Ming Li, Hangyuan Du, Yue Wang, and Edwin Hancock. Exploring the over-smoothing problem of graph neural networks for graph classification: An entropy-based viewpoint. In *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence (IJCAI)*, pages 3235–3243, 2025.

Siyi Qian, Haochao Ying, Renjun Hu, Jingbo Zhou, Jintai Chen, Danny Z. Chen, and Jian Wu. Robust training of graph neural networks via noise governance. In *Proceedings of the Sixteenth ACM International Conference on Web Search and Data Mining*, WSDM '23, page 607–615, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9781450394079. doi: 10.1145/3539597.3570369. URL `https://doi.org/10.1145/3539597.3570369`.

Scott E. Reed, Honglak Lee, Dragomir Anguelov, Christian Szegedy, Dumitru Erhan, and Andrew Rabinovich. Training deep neural networks on noisy labels with bootstrapping. In *ICLR 2015*, 2015. URL `http://arxiv.org/abs/1412.6596`.

T. Konstantin Rusch, Michael M. Bronstein, and Siddhartha Mishra. A survey on oversmoothing in graph neural networks. *ArXiv*, abs/2303.10993, 2023. URL `https://api.semanticscholar.org/CorpusID:257632346`.

Damian Szklarczyk, Annika Gable, David Lyon, Alexander Junge, Stefan Wyder, Jaime Huerta-Cepas, Milan Simonovic, Nadezhda Doncheva, John Morris, Peer Bork, Lars Jensen, and Christian von Mering. String v11: protein-protein association networks with increased coverage, supporting functional discovery in genome-wide experimental datasets. *Nucleic acids research*, 47, 11 2018a. doi: 10.1093/nar/gky1131.

Damian Szklarczyk, Annika L Gable, David Lyon, Alexander Junge, Stefan Wyder, Jaime Huerta-Cepas, Milan Simonovic, Nadezhda T Doncheva, John H Morris, Peer Bork, Lars J Jensen, and Christianvon Mering. STRING v11: protein–protein association networks with increased coverage, supporting functional discovery in genome-wide experimental datasets. *Nucleic Acids Research*, 47(D1):D607–D613, 11 2018b. ISSN 0305-1048. doi: 10.1093/nar/gky1131. URL `https://doi.org/10.1093/nar/gky1131`.

Nikil Wale and George Karypis. Comparison of descriptor spaces for chemical compound retrieval and classification. In *Sixth International Conference on Data Mining (ICDM'06)*, pages 678–689, 2006. doi: 10.1109/ICDM.2006.39.

Farooq Ahmad Wani, Maria Sofia Bucarelli, and Fabrizio Silvestri. Learning with noisy labels through learnable weighting and centroid similarity. *2024 International Joint Conference on Neural Networks (IJCNN)*, pages 1–9, 2023. URL `https://api.semanticscholar.org/CorpusID:257557810`.

Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural networks? In *7th International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*. OpenReview.net, 2019. URL `https://openreview.net/forum?id=ryGs6iA5Km`.

Shilong Xu, Jipeng Guo, Jiahao Long, Mingliang Cui, and Youqing Wang. High-low frequency filtering aware graph neural networks with representation decoupling learning. In *2024 China Automation Congress (CAC)*, pages 4005–4010, 2024. doi: 10.1109/CAC63892.2024.10864603.

Yujun Yan, Milad Hashemi, Kevin Swersky, Yaoqing Yang, and Danai Koutra. Two sides of the same coin: Heterophily and oversmoothing in graph convolutional neural networks. In *2022 IEEE International Conference on Data Mining (ICDM)*, pages 1287–1292, 2022. doi: 10.1109/ICDM54844.2022.00169.

Pinar Yanardag and S.V.N. Vishwanathan. Deep graph kernels. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '15, page 1365–1374, New York, NY, USA, 2015a. Association for Computing Machinery. ISBN 9781450336642. doi: 10.1145/2783258.2783417. URL `https://doi.org/10.1145/2783258.2783417`.

Pinar Yanardag and S.V.N. Vishwanathan. Deep graph kernels. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '15, page 1365–1374, New York, NY, USA, 2015b. Association for Computing Machinery. ISBN 9781450336642. doi: 10.1145/2783258.2783417. URL `https://doi.org/10.1145/2783258.2783417`.

Nan Yin, Li Shen, Mengzhu Wang, Xiao Luo, Zhigang Luo, and Dacheng Tao. Omg: Towards effective graph classification against label noise. *IEEE Transactions on Knowledge and Data Engineering*, 35(12):12873–12886, 2023. doi: 10.1109/TKDE.2023.3271677.

Jingyang Yuan, Xiao Luo, Yifang Qin, Yusheng Zhao, Wei Ju, and Ming Zhang. Learning on graphs under label noise. In *ICASSP 2023 - 2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 1–5, 2023a. doi: 10.1109/ICASSP49357.2023.10096088.

Jinliang Yuan, Hualei Yu, Meng Cao, Jianqing Song, Junyuan Xie, and Chongjun Wang. Self-supervised robust graph neural networks against noisy graphs and noisy labels. *Applied Intelligence*, 53(21):25154–25170, aug 2023b. ISSN 0924-669X. doi: 10.1007/s10489-023-04836-6. URL `https://doi.org/10.1007/s10489-023-04836-6`.

Chiyuan Zhang, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. Understanding deep learning (still) requires rethinking generalization. *Commun. ACM*, 64(3):107–115, February 2021. ISSN 0001-0782. doi: 10.1145/3446776. URL `https://doi.org/10.1145/3446776`.

Hongyi Zhang, Moustapha Ciss'e, Yann N. Dauphin, and David Lopez-Paz. mixup: Beyond empirical risk minimization. *CoRR*, abs/1710.09412, 2017.

Dengyong Zhou and Bernhard Schölkopf. Regularization on discrete spaces. In Walter G. Kropatsch, Robert Sablatnig, and Allan Hanbury, editors, *Pattern Recognition*, pages 361–368, Berlin, Heidelberg, 2005. Springer Berlin Heidelberg. ISBN 978-3-540-31942-9.

Kaixiong Zhou, Xiao Huang, Daochen Zha, Rui Chen, Li Li, Soo-Hyun Choi, and Xia Hu. Dirichlet energy constrained learning for deep graph neural networks. *Advances in neural information processing systems*, 34:21834–21846, 2021.

A Additional Related Work

A.1 Learning under Label Noise

Beyond robust losses, additional families include sample relabelling [Arazo et al., 2019, Reed et al., 2015], two-network co-teaching and divide-and-mix approaches [Han et al., 2018, Li et al., 2020, Kim et al., 2023], and regularization through MixUp [Zhang et al., 2017]. Reweighting methods adapt sample importance via training dynamics or held-out probes [Liu and Tao, 2015, Pleiss et al., 2020]. Most of these were developed for image classification; their direct transfer to graph classification is non-obvious because graphs lack the strong feature regularities of images.

A.2 Graph Learning under Noise

Most graph-related noise work targets node classification, exploiting either homophily (RS-GNN [Dai et al., 2022], NRGNN [Dai et al., 2021], RTGNN [Yuan et al., 2023a], pairwise interactions [Du et al., 2023]) or self-supervision (pseudo-label denoising [Yuan et al., 2023b, Qian et al., 2023]). Edge/feature noise has also been studied, e.g. structural noise [Fox and Rajamanickam, 2019, Dai et al., 2022]. These methods typically assume connected nodes share labels, which is incompatible with our graph-classification setting where labels are graph-level. The seminal work of NT et al. [2019] treats graph classification under label noise via a surrogate loss; OMG [Yin et al., 2023] combines contrastive learning, MixUp, and curriculum filtering. Both are practical improvements but provide no spectral or energy-based explanation of *why* GNNs fit graph-level noise.

A.3 Dirichlet Energy and Spectral Properties of GNNs

NT and Maehara [2019] and Rusch et al. [2023] analyze GNNs as low-pass filters, with self-loops further shrinking eigenvalues. Cai and Wang [2020], Oono and Suzuki [2020] prove that GNN Dirichlet energy decays exponentially when the product of the largest singular value of the weight matrix and the largest Laplacian eigenvalue is below 1. Giovanni et al. [2023] disentangle smoothing/sharpening via positive/negative weight-eigenvalues. Zhou et al. [2021] (NeurIPS) uses energy regularization to prevent over-smoothing on *node* classification. Our setting is the opposite regime: too *much* energy under graph-level label noise.

Qian et al. [2025] introduce a structural-entropy viewpoint on over-smoothing for graph classification, showing that high-entropy regions are most damaged by over-smoothing. We invert and combine this with Proposition 1 to get Theorem 1: the graphs that suffer most from over-*sharpening* (the noise-fitting regime) are also the high-entropy graphs that carry the most discriminative signal. The two papers therefore identify symmetric ends of the same spectral failure.

A.4 Lipschitz/Stability Analyses

Chuang and Jegelka [2022] bound GIN outputs via the Tree Mover's Distance, and Davidson and Dym [2025] extend uniform-Lipschitz analyses to MPNNs through Hölder continuity of aggregation/combination/readout. Juvina et al. [2024] derive optimal Lipschitz constants $\vartheta = \phi(\lambda_K)$ for graph convolutions under non-negative weights and ReLU. Gama et al. [2019] show output stability under graph-shift-operator perturbations. These results characterize stability w.r.t. inputs; we instead study stability against *label* noise, mediated by the within-graph energy of learned representations.

B Proofs and Extended Theory

B.1 Proof of Proposition 1

Let $\Lambda^i = \{\lambda_u^i\}_{u=1}^{N^i}$ be the spectrum of the per-graph Laplacian $\mathbf{\Delta}^i$. Define the block-diagonal Laplacian

$$\bar{\mathbf{\Delta}} = \mathrm{blkdiag}(\mathbf{\Delta}^1, \dots, \mathbf{\Delta}^n), \qquad \det(\bar{\mathbf{\Delta}} - \lambda \mathbf{I}_{N_{\text{tot}}}) = \prod_{i=1}^{n} \det(\mathbf{\Delta}^i - \lambda \mathbf{I}_{N^i}), \tag{9}$$

so $\bar{\Lambda} = \bigcup_i \Lambda^i$. Let $\boldsymbol{\psi}_j^i$ be the j-th eigenvector of $\mathbf{\Delta}^i$ at eigenvalue λ_j^i. The vector

$$\bar{\boldsymbol{\psi}}_u = \big(0, \dots, 0, \boldsymbol{\psi}_j^i, 0, \dots, 0\big)^\top \in \mathbb{R}^{N_{\text{tot}}} \tag{10}$$

satisfies $\bar{\mathbf{\Delta}} \bar{\boldsymbol{\psi}}_u = \lambda_j^i \bar{\boldsymbol{\psi}}_u$. With $\bar{\mathbf{Z}} = [\mathbf{Z}^1 \| \dots \| \mathbf{Z}^n]$, the Dirichlet energy of the block-diagonal graph admits the spectral form

$$E^{\text{dir}}(\bar{\mathbf{Z}}) = \sum_{r=1}^{m'} \sum_{u=1}^{N_{\text{tot}}} \bar{\lambda}_u (\bar{\boldsymbol{\psi}}_u^\top \bar{\mathbf{Z}}_{:,r})^2 = \sum_i \sum_{r,j} \lambda_j^i (\boldsymbol{\psi}_j^{i\top} \mathbf{Z}_{:,r}^i)^2 = \sum_i E^{\text{dir}}(\mathbf{Z}^i), \tag{11}$$

hence $E^{\text{dir}}_{\text{set}}(\mathcal{D}) = \frac{1}{|\mathcal{D}|} \sum_i E^{\text{dir}}(\mathbf{Z}^i) = \frac{1}{|\mathcal{D}|} E^{\text{dir}}(\bar{\mathbf{Z}})$. □

B.2 Proof of Theorem 1

Step 1: Sharpening lifts spectral mass. Let $\Delta E^{\text{dir}}(\mathbf{Z}^i)$ denote the increase in $E^{\text{dir}}(\mathbf{Z}^i)$ across a training step that increases $E^{\text{dir}}_{\text{set}}$. By Proposition 1, the increase in $E^{\text{dir}}(\bar{\mathbf{Z}})$ is concentrated on eigenpairs $(\bar{\lambda}_u, \bar{\boldsymbol{\psi}}_u)$ with large $\bar{\lambda}_u$ (since the spectral coefficient is multiplied by $\bar{\lambda}_u$). Each such eigenpair lies entirely on a single connected component of $\bar{\mathcal{G}}$, i.e. a single graph, by the block-diagonal structure of $\bar{\mathbf{\Delta}}$.

Step 2: High-entropy graphs absorb more spectral mass. By Theorems 1–2 of Qian et al. [2025], for graphs with the same average sparsity, increasing the number of nodes (Theorem 1) or the connection density (Theorem 2) strictly increases the structural Shannon entropy. A higher-entropy graph also has a richer (more spread) Laplacian spectrum: its tail of large λ_j^i contains more mass than that of a lower-entropy graph at the same scale. Sharpening therefore allocates a larger share of $\Delta E^{\text{dir}}(\bar{\mathbf{Z}})$ to the high-entropy graph $\mathcal{G}^j$ than to the low-entropy graph $\mathcal{G}^i$. Formally, if $\sigma_i = \sum_{j:\, \lambda_j^i \geq \tau} \lambda_j^i$ is the high-frequency spectral mass of graph i above threshold τ, then $H(\mathcal{G}^j) \geq H(\mathcal{G}^i) \Rightarrow \sigma_j \geq \sigma_i$, and a sharpening update $\mathbf{Z}^i \to \mathbf{Z}^i + \delta$ aligned with the high-frequency subspace increases $E^{\text{dir}}(\mathbf{Z}^j)$ by at least as much as $E^{\text{dir}}(\mathbf{Z}^i)$.

Step 3: Pooling-stage discriminability collapses faster for high-entropy graphs. The classification head sees pooled features $\hat{f}(\mathbf{Z}^i) = \mathrm{POOL}(\mathbf{Z}^i)$. The pooled JSD between two graphs with similar pooled distributions P^i, P^j is bounded above by a function of within-graph node-distribution distinctness (Qian et al., 2025, Sec. 3.3). When sharpening of $\mathbf{Z}^j$ shrinks the per-node distribution differences within $\mathcal{G}^j$ (the over-smoothing-in-high-entropy-region phenomenon of Qian et al., 2025), the pooled descriptor of $\mathcal{G}^j$ approaches that of any other graph with similar high-entropy structure, including a mislabeled one. The model then has no choice but to bias the readout toward labels consistent with the corrupted training graph, completing the noise-fitting loop. The collapse is faster for $\mathcal{G}^j$ than $\mathcal{G}^i$ because $\sigma_j \geq \sigma_i$. $\square$

B.3 Why a Decrease of $E^{\mathrm{dir}}_{\mathrm{set}}$ Is Non-Trivial in Graph Classification

Three consecutive observations together justify the term "non-trivial":

(i) Noise is injected at the graph-label level, not at the node-feature level. Therefore, the increase of $E^{\mathrm{dir}}(\mathbf{Z}^i)$ during noise fitting cannot be a re-encoding of the noise itself; it must be a behavioral response of the network through training dynamics.

(ii) For a randomly-initialized network, $E^{\mathrm{dir}}_{\mathrm{set}}$ is independent of class labels by construction. Any correlation between $E^{\mathrm{dir}}_{\mathrm{set}}$ and labels emerges from the optimization process, not from the data.

(iii) By Theorem 1, the only optimization path that fits incorrect graph-level labels is to lift spectral mass into high-entropy directions of $\bar{\mathbf{\Delta}}$, which are precisely the directions that a clean, generalization-driven network should preserve. Constraining $E^{\mathrm{dir}}_{\mathrm{set}}$ therefore selectively blocks the noise-fitting path while allowing the clean-learning path, an alignment that does *not* hold for node classification, where label noise corrupts the same nodes whose smoothness the energy measures, and where decreasing energy can erase the discriminative signal.

These three points collectively make the connection between energy reduction and noise robustness in graph classification a non-trivial property. The analogous statement in node classification reduces to "the model fits homophily, and homophily is what we measured."

C Experimental Setting

C.1 Datasets

ogbg-ppa [Szklarczyk et al., 2018a]: 158,100 protein-association graphs across 37 taxonomic groups, derived from STRING [Szklarczyk et al., 2018b]; the standard OGB graph-property-prediction benchmark. Modified PPA-30%-6 takes 30% of training graphs from 6 selected classes, used as a controlled testbed. **PROTEINS** [Borgwardt et al., 2005]: 1,113 protein graphs; 2 classes; mean 39 nodes/graph. **ENZYMES** [Borgwardt et al., 2005]: 600 enzyme graphs across 6 EC top-level classes. **MSRC-21** [Neumann et al., 2016]: 563 image-segmentation graphs, 20 classes. **MUTAG** [Kriege and Mutzel, 2012]: 188 molecular graphs, 2 classes. **IMDB-BINARY** [Yanardag and Vishwanathan, 2015b]: 1,000 social graphs, 2 classes. **REDDIT-MULTI-12K** [Yanardag and Vishwanathan, 2015b]: 11,929 thread-graphs, 11 classes. **MNIST-graph**: 55,000 superpixel graphs over the MNIST images, 10 classes.

C.2 Synthetic Datasets for Failure-Mode F2

We sample average graph orders $\bar{N} \in \{5, 10, \ldots, 60\}$, drawing actual node counts from $\mathrm{Poisson}(\bar{N})$. Each dataset has 6 classes with 1,400 graphs/class, fixed average degree 2, edges sampled uniformly. Node features are $\mathcal{N}(\mu_c, 1.5)$ with $\mu_c = f(\text{class})$. Symmetric noise rate is fixed at 35%.

C.3 Architectures and Hyperparameters

Unless otherwise stated, we use a 5-layer GIN [Xu et al., 2019] or GCN [Kipf and Welling, 2017] with 300 hidden units, Adam ($\eta = 10^{-3}$, weight decay 10^{-4}), batch size 32, 200–1000 epochs, $\eta_u = 1$ for GCOD's outlier parameter. Compute: NVIDIA RTX A6000.

C.4 Spectral Bias Implementation Details (CE+W2)

For each layer l, after a gradient update, decompose $\mathbf{W}_l^2 = \mathbf{\Phi}_l^2 \boldsymbol{\mu}_l^2 (\mathbf{\Phi}_l^2)^{-1}$ and replace by $\mathbf{\Phi}_l^2 [\boldsymbol{\mu}_l^2]^+ (\mathbf{\Phi}_l^2)^{-1}$, where $[\cdot]^+$ is element-wise ReLU on the eigenvalues. The projection is detached from the gradient graph. We apply it to the post-aggregation matrix $\mathbf{W}^2$ of every layer; applying it also to $\mathbf{W}^1$ marginally improves smoothing but degrades stability. Convergence variance under this scheme is shown in Fig. 8.

C.5 Class-Specific and Fixed Bounds for $\mathcal{L}_{DE}$

Class-specific: after every epoch, set $U_c = |\mathcal{D}_c^{\text{val}}|^{-1} \sum_{i\in\mathcal{D}_c^{\text{val}}} E^{\text{dir}}(\mathbf{Z}^i)$. The validation set must be clean to keep U_c class-specific. **Fixed:** $U_c = U$ globally. We progressively decrease U from a large initial value until accuracy under clean labels begins to drop; the final U is the largest value that achieves a noticeable energy ceiling without over-smoothing.

C.6 Design of GCOD

We adopt the centroid-based soft labels of Wani et al. [2023]: normalize $\phi(\mathbf{z}_i) = \mathbf{z}_i/\|\mathbf{z}_i\|$; per-class centroid $\boldsymbol{\phi}_c = \bar{\mathbf{z}}_c/\|\bar{\mathbf{z}}_c\|$; soft target $\tilde{\mathbf{y}}_i = [\max(\phi(\mathbf{z}_i)\boldsymbol{\phi}_c^\top, 0)\mathbf{y}_i]$. Equations (6)–(8) are then optimized as in the main text. The training-accuracy feedback in $\mathcal{L}_1$ and $\mathcal{L}_3$ is critical: without it, $\mathbf{u}$ dominates early in training and discounts samples before genuine patterns are learned (Fig. 7).

D Failure-Mode Experiments (Extended)

Additional plots for the F1–F3 failure modes are provided here for completeness, including effect-of-dataset-size for both PPA and synthetic datasets.

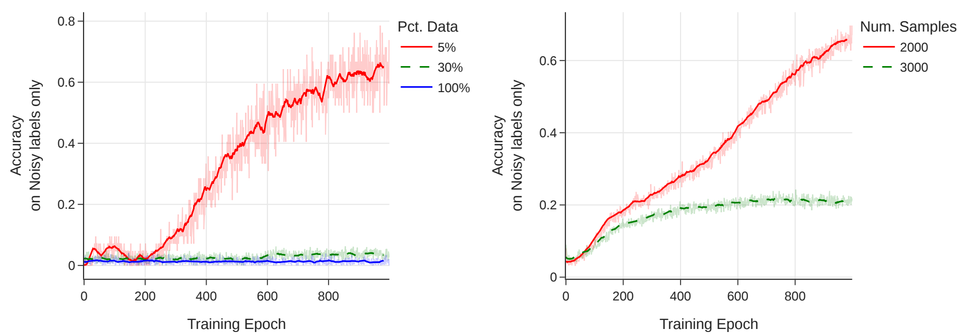

(a) (b)

Figure 3: Effect of training-set size on noise memorization (training accuracy on noisy labels only). (a) PPA subsampling at 20% symmetric noise. (b) Synthetic graphs of order 7, 6 classes, 35% noise.

E Dirichlet-Energy Amplification in High-Frequency Components

We use the HLFF-GNN framework [Xu et al., 2024] (implemented as `FGRLConv`) to decompose representations into low-frequency Z_1, high-frequency Z_2, and residual Y. Under the composite loss

$$\mathcal{L} = \|Y - X\|_F^2 + \lambda\,\mathrm{tr}(Z_1^\top \mathbf{\Delta} Z_1) + \alpha\,\mathrm{tr}(Z_2^\top(\mathbf{I}-\mathbf{\Delta})Z_2) + \beta(\|Y^\top Z_1\|_F^2 + \|Y^\top Z_2\|_F^2), \quad (12)$$

the spectral form of the high-frequency energy is $E^{\text{dir}}(Z_2) = \sum_{r,u} \lambda_u(\boldsymbol{\psi}_u^\top Z_{2,r})^2$, in which large λ_u amplify any noise-driven content. On ENZYMES with 30% symmetric noise, we observe a sharp monotone increase in $E^{\text{dir}}(Z_2)$, while $E^{\text{dir}}(Z_1)$ and $E^{\text{dir}}(Y)$ remain stable, consistent with Theorem 1.

F Additional Empirical Studies on Energy Behavior

Energy under clean overfitting vs. noise. On ENZYMES without regularization, a deliberately overfit GIN exhibits rising training accuracy and rising $E^{\text{dir}}_{\text{set}}$, confirming that energy tracks overfitting in general. However, the rise under noise is markedly steeper.

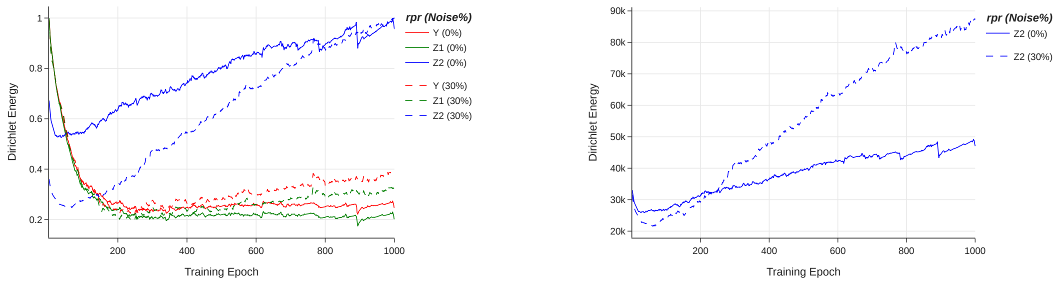

(a) Normalized $E^{\text{dir}}(Y, Z_1, Z_2)$.

(b) $E^{\text{dir}}(Z_2)$ alone.

Figure 4: HLFF-GNN frequency decomposition on ENZYMES. Solid lines: clean labels. Dashed: 30% symmetric noise. The high-frequency component Z_2 is the only one whose energy diverges under noise, supporting the theoretical prediction.

Energy scales with noise rate. On the controlled PPA subset, we sweep symmetric noise $\in \{10, 20, 30, 40, 60\}\%$. All curves follow the same U-shape (initial decrease, then rise), and the asymptotic energy increases monotonically with noise rate. This monotonicity is what makes $E^{\text{dir}}_{\text{set}}$ a diagnostic.

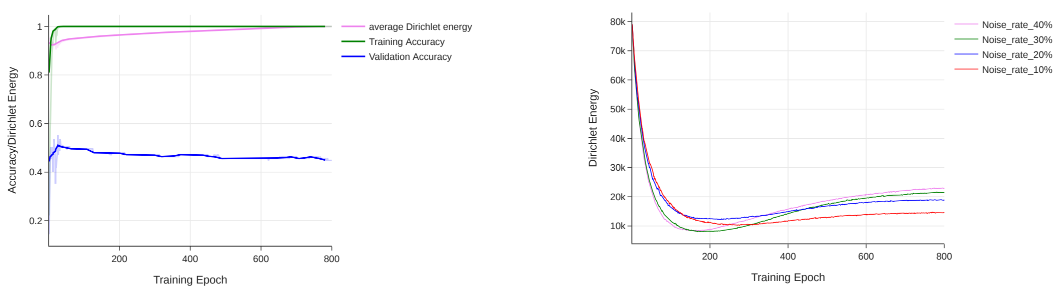

(a) Clean overfitting on ENZYMES.

(b) $E^{\text{dir}}_{\text{set}}$ vs. noise rate on PPA.

Figure 5: (a) Even without label noise, energy rises during overfitting; the surge under noise is much steeper. (b) Higher noise rates yield monotonically higher final energy.

G Additional Tables and Figures

Table 5: Sensitivity of GCOD to the learning rate η_u of the outlier parameter $\mathbf{u}$, under 40% asymmetric noise.

Dataset	$\eta_u = 1$	$\eta_u = 0.1$	% Change
PROTEINS	76.19	75.89	-0.39%
MNIST	72.61	72.48	-0.18%
ENZYMES	69.81	68.33	-2.12%
IMDB-B	72.89	71.94	-1.30%
MUTAG	93.19	92.61	-0.62%
REDDIT	47.94	48.08	$+0.29\%$
MSRC-21	95.57	95.45	-0.13%

Table 6: Performance of CE vs. GCOD with 20% symmetric label noise. Last column reports the absolute gain (GCOD−CE) on test accuracy.

Dataset	Metric	0% CE	20% CE	20% GCOD	20% (GCOD−CE)
MNIST	Best	72.69	66.30	71.26	+4.96
	Last	69.80	53.64	67.88	+14.24
	Diff	2.89	12.66	3.38	
ENZYMES	Best	73.33	64.16	68.69	+4.53
	Last	65.78	57.50	62.54	+5.04
	Diff	7.55	6.66	6.15	
MSRC-21	Best	96.69	90.26	94.69	+4.43
	Last	93.80	79.64	90.26	+10.62
	Diff	2.89	10.62	4.43	
PROTEINS	Best	81.16	76.23	79.38	+3.15
	Last	79.18	62.32	78.12	+15.80
	Diff	1.98	13.91	1.26	
MUTAG	Best	94.73	89.47	90.01	+0.54
	Last	84.21	68.42	86.84	+18.42
	Diff	10.52	21.05	3.17	
IMDB-B	Best	76.50	75.00	75.40	+0.40
	Last	71.60	71.00	73.50	+2.50
	Diff	4.90	4.00	1.90	
REDDIT	Best	48.15	45.05	44.98	−0.07
	Last	46.01	44.67	44.89	+0.22
	Diff	2.14	0.38	0.09	

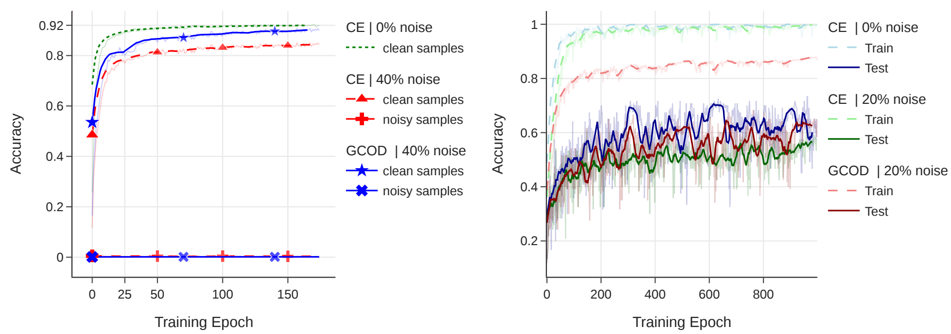

(a) (b)

Figure 6: (a) Train accuracy on known noisy and clean PPA subsets (4-class, 40% symmetric noise) for GCN with CE: the model never separates noisy from clean training samples without a robust loss. (b) Train/test accuracy on ENZYMES with GIN under clean and 20% symmetric noise.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The abstract and Section 1 state precisely the claims supported by the theoretical results in Section 6 and the experiments in Section 8.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: A dedicated paragraph in the Conclusion discusses the absence of a closed-form noise-energy characterization, reliance on synthetic noise, and the lack of standardized real-world noisy graph-classification benchmarks.

3. **Theory assumptions and proofs**

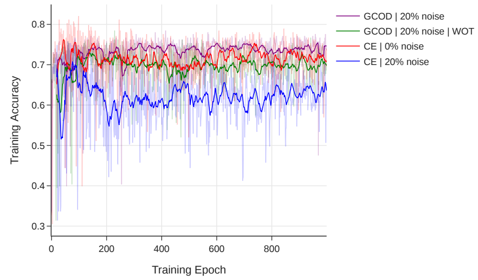

Figure 7: Ablation on the training-accuracy feedback in GCOD. Without scaling u by a_{train} ("GCOD WOT"), the outlier parameter dominates early and harms generalization; the full GCOD activates outlier-discounting gradually.

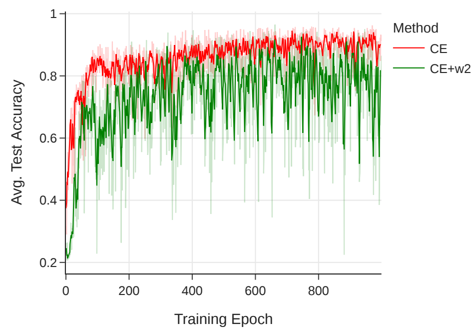

Figure 8: Test accuracy on PPA-30%-6 for CE and CE+W2. CE+W2 reaches comparable peak accuracy but exhibits unstable convergence due to the per-epoch eigendecomposition.

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Proposition 1 and Theorem 1 state assumptions explicitly in the main text, and full proofs are provided in Appendix B.

4. **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper?

Answer: [Yes]

Justification: Appendix C provides datasets, splits, hyperparameters, architectures, and the implementation details of all three methods. Anonymized code is provided in the supplemental material.

5. **Open access to data and code**

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results?

Answer: [Yes]

Justification: All datasets are public benchmarks; an anonymous code repository accompanies the submission.

6. **Experimental setting/details**

Question: Does the paper specify all the training and test details necessary to understand the results?

Answer: [Yes]

Justification: Section 8 and Appendix C report all training and evaluation details.

7. **Experiment statistical significance**

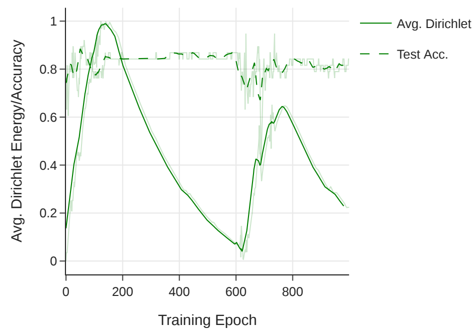

Figure 9: Average test accuracy and average within-graph $E_{\text{set}}^{\text{dir}}$ on MUTAG (0% noise) with GIN: classification accuracy improves while energy decays, the clean-data signature of a healthy GNN.

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All main results report mean$\pm$std over 4 independent runs with different seeds.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources?

   Answer: [Yes]

   Justification: Appendix C reports the use of NVIDIA RTX A6000 GPUs and Table 4 reports relative training-time overheads.

9. **Code of ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics?

   Answer: [Yes]

   Justification: The work uses public benchmarks, does not involve human subjects, and the broader-impact paragraph discusses application contexts.

10. **Broader impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: [Yes]

    Justification: Discussed in the Conclusion. The work primarily benefits scientific applications (chemistry, biology) where label noise is systematic.

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse?

    Answer: [N/A]

    Justification: The paper does not release new datasets or pretrained models with misuse risk.

12. **Licenses for existing assets**

    Question: Are the creators or original owners of assets, properly credited and are the license and terms of use explicitly mentioned and properly respected?

    Answer: [Yes]

    Justification: All datasets are cited; OGB, TU-Datasets and PyTorch-Geometric implementations are used under their respective licenses.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The released code provides documentation for the proposed methods.

14. **Crowdsourcing and research with human subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [N/A]

    Justification: No human-subject or crowdsourcing experiments.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [N/A]

    Justification: No human subjects.

16. **Declaration of LLM usage**

    Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research?

    Answer: [N/A]

    Justification: LLMs are not part of the core methods.