
Energy-Guided Smoothness for Robust Graph Classification under Label Noise

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Graph Neural Networks (GNNs) are vulnerable to label noise in graph classification, yet the mechanisms driving this vulnerability are poorly understood. We
2 empirically identify a property of the *learned representations* that tracks noise
3 memorization: the within-graph Dirichlet energy of node features. Energy decreases
4 as the model captures genuine class structure, then rises sharply, mainly in
5 high-frequency components, when the model begins to memorize corrupted graph
6 labels. We support this observation with a Poincaré–bi-Lipschitz lower bound on
7 the energy implied by noise fitting, combined with an entropy-based asymmetry
8 result adapted from Qian et al. [2025] indicating that the energy growth is preferentially
9 absorbed by graphs most informative for classification. Guided by these
10 observations we evaluate three interventions: a Dirichlet-energy regularizer ($\mathcal{L}_{DE}$),
11 a robust loss (GCOD), and a spectral weight constraint (CE+W2). Across eight
12 benchmarks, including the full and modified ogbg-ppa, symmetric noise up to 60%,
13 asymmetric noise, and a transfer to Cora node classification, the three interventions
14 all reduce within-graph energy and improve robustness while preserving clean-data
15 performance. GCOD, our practical recommendation, matches OMG [Yin et al.,
16 2023] at $\sim 70\times$ lower training overhead.
17

18 1 Introduction

19 Graph Neural Networks (GNNs) are now the workhorse of *graph classification*, the task of assigning
20 a single label to an entire graph, with applications spanning molecular property prediction [Wale and
21 Karypis, 2006], bioinformatics [Borgwardt et al., 2005], and social-network analysis [Yanardag and
22 Vishwanathan, 2015a]. In real deployments, the labels attached to graphs are routinely noisy: assays
23 disagree, annotators err, and weak supervision injects systematic mistakes. Although learning under
24 label noise has been thoroughly studied for images [Algan and Ulusoy, 2021] and for node classification
25 [Dai et al., 2021, Yuan et al., 2023a], graph classification under noisy labels is dramatically less
26 explored [NT et al., 2019, Yin et al., 2023].

27 **Graph classification is a fundamentally different setting.** In node classification, label noise
28 corrupts *individual node labels* on a single graph, and homophily, the tendency of connected nodes
29 to share labels, is a structural prior that directly couples Dirichlet energy of the *label signal* to
30 noise. In graph classification, label noise flips the label of an *entire graph*: nodes have no individual
31 classification labels, and within-graph homophily/heterophily is largely orthogonal to the graph-level
32 label. The standard tools and intuitions from node classification therefore do not transfer.

33 **Our central observation.** We measure the Dirichlet energy of *learned node representations*, not of
34 labels. Empirically, this energy first decreases as the model captures genuine, transferable patterns;

then, when the model begins to memorize incorrect graph-level labels, it rises sharply, with the surge concentrated in the high-frequency spectrum. This rise is not a tautology: noise is injected only at the graph-label level, while node features and edges are unchanged, so any correlation between energy and labels has to emerge through training dynamics. Theorem 1 supports this observation analytically: under a bi-Lipschitz regularity assumption, fitting noisy graph-level labels for two structurally similar graphs forces a lower bound on their within-graph Dirichlet energies via the Graph Poincaré inequality.

Why does suppressing this energy growth help? Combining a dataset-level spectral identity (Proposition 1) with an entropy argument from Qian et al. [2025], Section 6 argues that uncontrolled energy growth is preferentially absorbed by *high-entropy graphs*, the ones carrying most discriminative signal at pooling. Reducing within-graph Dirichlet energy therefore plausibly suppresses noise fitting while preserving the high-entropy regions classification accuracy depends on, a hypothesis we test with three mechanistically distinct interventions.

Contributions.

- A controlled study of *when* GNNs fit graph-level label noise (Section 4), isolating three failure modes: over-parameterization, low-order graphs, and small training sets.
- Within-graph Dirichlet energy as a diagnostic for graph-level noise memorization (Section 5), with a frequency decomposition that localizes the surge to high-frequency components and an empirical scaling with the noise rate.
- Analytical support for the diagnostic (Section 6): a spectral identity (Proposition 1), a Poincaré–bi-Lipschitz lower bound under noise fitting (Theorem 1), and an entropy-based asymmetry result (Corollary 1). These results are unique to graph classification, where labels and feature smoothness live on different objects, unlike node classification.
- Three interventions (Section 7): a Dirichlet-energy regularizer $\mathcal{L}_{DE}$, a centroid-based loss GCOD, and a spectral weight constraint CE+W2. They are mechanistically distinct and all reduce within-graph energy while improving robustness; GCOD is our practical recommendation.
- Evaluation on eight benchmarks under symmetric, asymmetric, and high-rate (60%) noise, the full and a modified ogbg-ppa, and a transfer to Cora node classification. GCOD matches OMG [Yin et al., 2023] and SOP [Liu et al., 2022] at $\sim 70\times$ lower compute.

2 Related Work

Learning under label noise. Robust loss functions [Ghosh et al., 2017, Patrini et al., 2017], sample re-weighting [Liu and Tao, 2015, Liu et al., 2022], early-learning regularizers [Arpit et al., 2017], and small-loss filtering [Gui et al., 2021] dominate the image-classification literature. Wani et al. [2023] introduce NCOD, a centroid-based outlier-discounting loss that we adapt to graphs.

Noisy-label learning on graphs. Most prior work concerns *node* classification, exploiting homophily, contrastive losses, or pseudo-labels [Dai et al., 2021, Yuan et al., 2023a,b, Li et al., 2024, Kang et al., 2018]. For *graph* classification, NT et al. [2019] propose a surrogate loss that discards estimated-noisy labels under specific assumptions, and OMG [Yin et al., 2023] combines contrastive learning, MixUp [Lim et al., 2022], and curriculum filtering. We differ in two ways: (i) we provide a spectral/energy-based explanation of *why* GNNs fit graph-level noise, rather than only mitigating it, and (ii) our practical method (GCOD) achieves better or comparable robustness at a fraction of the compute.

Dirichlet energy and over-smoothing. Dirichlet energy quantifies signal smoothness on graphs [Zhou and Schölkopf, 2005]. Existing analyses establish that GNNs act as low-pass filters whose energy decays exponentially with depth under spectral conditions [NT and Maehara, 2019, Cai and Wang, 2020, Oono and Suzuki, 2020, Rusch et al., 2023], and that weight-matrix eigenvalues control attraction/repulsion of node features [Giovanni et al., 2023]. Energy-based regularizers [Zhou et al., 2021, Bo et al., 2021, Chen et al., 2023] address *too little* energy (over-smoothing) on *node* classification. We address the opposite regime: *too much* within-graph energy caused by graph-level label noise. Qian et al. [2025] recently show that, for graph classification specifically, over-smoothing in *high-entropy* regions is the most damaging; we leverage this insight (Corollary 1) to argue why suppressing energy growth is essential under graph-level noise. A more comprehensive related-work discussion is in Appendix A.

3 Background

Notation. A graph is $\mathcal{G} = (\mathcal{V}, \mathcal{E}, \mathbf{X})$ with $N = |\mathcal{V}|$, degree matrix $\mathbf{D}$, adjacency $\mathbf{A} \in \{0, 1\}^{N \times N}$, feature matrix $\mathbf{X} \in \mathbb{R}^{N \times m}$, combinatorial Laplacian $\mathbf{L} = \mathbf{D} - \mathbf{A}$ (with eigenvalues $0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_N$) and normalized Laplacian $\Delta = \mathbf{D}^{-1/2} \mathbf{L} \mathbf{D}^{-1/2}$. We write $\lambda_2(\mathcal{G})$ for $\lambda_2(\mathbf{L})$ throughout and use Δ only when explicitly stated. In *graph classification*, the dataset is $\mathcal{D} = \{(\mathcal{G}^i, \mathbf{y}_i)\}_{i=1}^n$, $\mathbf{y}_i \in \{0, 1\}^{|C|}$ the one-hot graph label; we write c_i for the class. Under message passing, $\mathbf{H}_i^{l+1} = \text{UP}_{\Omega_l}(\mathbf{H}_i^l, \text{AGGR}_{\mathbf{W}_l}(\mathbf{H}_i^l, \mathbf{A}^i))$, with $\mathbf{H}_i^0 \equiv \mathbf{X}^i$ and $\mathbf{Z}^i \equiv \mathbf{H}_i^L$. A permutation-invariant readout $f_\theta : \mathbb{R}^{N_i \times m'} \rightarrow \mathbb{R}^{|C|}$ maps node features to class probabilities $\hat{\mathbf{y}}_i$. Under noisy labels, the observed c_i may differ from the ground truth; the test set is clean.

Dirichlet energy on graphs. For graph $\mathcal{G}^i$ with representation $\mathbf{Z}^i \in \mathbb{R}^{N_i \times m'}$ we use the combinatorial Dirichlet energy

$$E^{\text{dir}}(\mathbf{Z}^i) = \sum_{(u,v) \in \mathcal{E}^i} \|\mathbf{Z}_u^i - \mathbf{Z}_v^i\|_2^2 = \text{tr}((\mathbf{Z}^i)^\top \mathbf{L}^i \mathbf{Z}^i). \quad (1)$$

Small E^{dir} means smooth (low-frequency) representations within the graph; large E^{dir} means sharpening (high-frequency). The symmetrically-normalized variant $E_{\text{sym}}^{\text{dir}}(\mathbf{Z}) = \sum_{(u,v)} \|\mathbf{Z}_u / \sqrt{d_u} - \mathbf{Z}_v / \sqrt{d_v}\|_2^2 = \text{tr}(\mathbf{Z}^\top \Delta \mathbf{Z})$, used in some over-smoothing analyses [NT and Maehara, 2019, Oono and Suzuki, 2020], is equivalent up to degree-bounded constants and our qualitative claims hold for either; we use (1) because it pairs cleanly with the Poincaré inequality of Section 6. To track behavior across a class-restricted subset $\mathcal{S}$ of training graphs we use

$$E_{\text{set}}^{\text{dir}}(\mathcal{S}) = \frac{1}{|\mathcal{S}|} \sum_{\mathcal{G}^i \in \mathcal{S}} E^{\text{dir}}(\mathbf{Z}^i). \quad (2)$$

Crucially, (1) is computed on *learned representations*, not on labels. Within-graph homophily of features therefore evolves with training, decoupled from the graph-level label of $\mathcal{G}^i$.

4 When GNNs Fit Graph-Label Noise

GNNs trained with standard cross-entropy (CE) on graph classification tasks are not uniformly vulnerable to label noise. On the full ogbg-ppa [Szklarczyk et al., 2018a] (37 classes, the standard OGB benchmark), even 40% symmetric noise barely degrades CE accuracy, because the model is under-parameterized for the full task and behaves robustly [Zhang et al., 2021]. To expose the noise-fitting regime, we control three axes: **(F1)** number of classes (proxy for over-parameterization), **(F2)** graph order (size of each graph), and **(F3)** training-set size. We use a 5-layer GIN [Xu et al., 2019] with 300 hidden units throughout, unless stated.

In all three regimes (Fig. 1 and Appendix D), GNNs progressively memorize noisy graph-level labels. The natural next question is therefore: what changes in the network’s internal representations during memorization, and can this change be detected and controlled? The next section identifies Dirichlet energy as that internal signal.

5 Dirichlet Energy as a Diagnostic for Noise Fitting

We track $E_{\text{set}}^{\text{dir}}(\mathcal{D}_c)$, the within-graph energy averaged over training graphs of class c . Across architectures and datasets (Fig. 2), the diagnostic is consistent: $E_{\text{set}}^{\text{dir}}$ falls during the early “learning” phase, then rises precisely when the model begins to fit noisy graphs (CE-20% in Fig. 2b), with 40% noise yielding a strictly larger end-of-training energy than 10% noise (Appendix F).

Energy is sensitive to noise specifically through high-frequency components. A natural concern is that energy also rises during clean-data overfitting [Yan et al., 2022]; we confirm this in Appendix F but show the noise-driven surge is markedly steeper. To localize the surge spectrally, we use the HLFF-GNN decomposition [Xu et al., 2024] to split representations into low-frequency (Z_1) and high-frequency (Z_2) components. On ENZYMES with 30% noise, $E^{\text{dir}}(Z_1)$ remains essentially unchanged, while $E^{\text{dir}}(Z_2)$ rises monotonically, and slope/variance statistics of $E^{\text{dir}}(Z_2)$ provide an early warning for label noise (Appendix E). HLFF-GNN is used purely as an analytical tool; the same total-energy surge appears in standard GIN/GCN.




Figure 1: Training accuracy on *noisy labels only*: a higher curve means more memorization. (a) Reducing the number of PPA classes (with 20% symmetric noise) makes a fixed GIN over-parameterized, accelerating noise fitting. (b) Smaller average graph order (synthetic data, 35% noise) leaves less internal structure for averaging, enabling memorization.

132 **Why energy is more useful than a generic overfitting signal.** Standard noise-detection signals
 133 require a *clean* held-out set: validation loss is uninformative when validation labels follow the same
 134 noise distribution as training (the noisy validation loss continues to decrease as the model fits noise,
 135 as we observe in our setting); train–test gap is unobservable at training time; small-loss heuristics
 136 need a calibration regime. Within-graph Dirichlet energy is computed entirely on the training graphs
 137 and their representations, so it remains a valid noise-fitting indicator under fully-noisy supervision.
 138 Beyond replacing the missing clean signal, it tells you *how* the network is failing: which spectral
 139 components drive the surge. This mechanistic specificity is what turns a diagnostic into a design
 140 principle, and motivates the three interventions in Section 7, each of which targets the same harmful
 141 spectral behavior in a distinct way.

142 6 Analytical Support: Why Energy Suppression Helps

143 Section 5 established Dirichlet energy as a reliable empirical signal. We now show this is not just
 144 correlation: under standard regularity, within-graph energy is *provably forced* to grow when a GNN
 145 fits graph-level label noise, with the growth absorbed disproportionately by structurally rich graphs.
 146 The argument combines three pieces: a dataset-level spectral identity (Proposition 1), the Graph
 147 Poincaré inequality (Lemma 1), and a bi-Lipschitz lower bound on the GNN (Lemma 2) [Chuang
 148 and Jegelka, 2022, Davidson and Dym, 2025, Juvina et al., 2024], yielding Theorem 1.

149 6.1 Three Building Blocks

150 **Proposition 1** (Dataset-level Dirichlet energy). *Let $\mathcal{D} = \{\mathcal{G}^1, \dots, \mathcal{G}^n\}$ with representations*
 151 *$\mathbf{Z}^i \in \mathbb{R}^{N^i \times m'}$ and adjacencies $\mathbf{A}^i$. Define the block-diagonal graph $\bar{\mathcal{G}} = (\bar{\mathbf{Z}}, \bar{\mathbf{A}})$ with $\bar{\mathbf{Z}} =$*
 152 *$[\mathbf{Z}^1 \parallel \dots \parallel \mathbf{Z}^n] \in \mathbb{R}^{N_{\text{tot}} \times m'}$ (vertical concatenation) and $\bar{\mathbf{A}} = \text{blkdiag}(\mathbf{A}^1, \dots, \mathbf{A}^n)$, with associated*
 153 *combinatorial Laplacian $\bar{\mathbf{L}} = \text{blkdiag}(\mathbf{L}^1, \dots, \mathbf{L}^n)$. Then*

$$E_{\text{set}}^{\text{dir}}(\mathcal{D}) = \frac{1}{|\mathcal{D}|} E^{\text{dir}}(\bar{\mathbf{Z}}) = \frac{1}{|\mathcal{D}|} \sum_{r=1}^{m'} \sum_{u=1}^{N_{\text{tot}}} \bar{\lambda}_u (\bar{\psi}_u^\top \bar{\mathbf{Z}}_{:,r})^2, \quad (3)$$

154 *where $\{(\bar{\lambda}_u, \bar{\psi}_u)\}$ are eigenpairs of $\bar{\mathbf{L}}$, and the spectrum decomposes as $\bar{\mathbf{L}} = \bigcup_i \Lambda^i$ with each*
 155 *eigenvector supported on a single connected component $\mathcal{G}^i$.*

156 The proof (Appendix B) follows from the block-diagonal structure of $\bar{\mathbf{L}}$. The key structural fact: *each*
 157 *high-frequency eigenvector of $\bar{\mathbf{L}}$ is owned by exactly one graph*, so sharpening along it raises only
 158 *that graph’s energy.*

159 **Lemma 1** (Graph Poincaré, vector-valued form). *For any feature matrix $\mathbf{Z} \in \mathbb{R}^{N \times m'}$ on a connected*
 160 *graph $\mathcal{G}$ with combinatorial Laplacian $\mathbf{L} = \mathbf{D} - \mathbf{A}$ and second-smallest eigenvalue $\lambda_2(\mathcal{G}) > 0$, let*




Figure 2: Within-graph Dirichlet energy of learned representations *rises during noise memorization*, across architectures (GIN/GCN) and datasets (PPA, MUTAG, MNIST, PROTEINS). Under clean labels, $E_{\text{set}}^{\text{dir}}$ steadily decays as accuracy improves; under noise, it follows the same decay until the model starts fitting noisy graphs, after which it sharply rises. GCOD suppresses this surge.

161 $\bar{\mathbf{z}} = \frac{1}{N} \sum_v \mathbf{Z}_v$. Then

$$\sum_{v=1}^N \|\mathbf{Z}_v - \bar{\mathbf{z}}\|_2^2 \leq \frac{1}{\lambda_2(\mathcal{G})} E^{\text{dir}}(\mathbf{Z}). \quad (4)$$

162 (4) bounds within-graph feature variance by Dirichlet energy scaled by the spectral gap; proof by
163 Courant–Fischer per feature dimension (Appendix B).

164 **Lemma 2** (Bi-Lipschitz lower bound for a GIN block with 2-layer combine, adapted from Chuang and
165 Jegelka, 2022, Davidson and Dym, 2025, Juvina et al., 2024). *Let $h_\theta(\mathbf{X}) = \sigma_2(\sigma_1(\mathbf{M}\mathbf{X}\mathbf{W}_1)\mathbf{W}_2)$
166 be a single GIN aggregation $\mathbf{M} = (1 + \varepsilon)\mathbf{I} + \mathbf{A} \in \mathbb{R}^{N \times N}$ followed by a two-layer MLP with
167 weights $\mathbf{W}_1 \in \mathbb{R}^{m \times m'}$ and $\mathbf{W}_2 \in \mathbb{R}^{m' \times m''}$ acting on the feature dimension, with the widening
168 regime $m \leq m' \leq m''$ (so the linear maps have full column rank). Assume $\sigma_{\min}(\mathbf{M}) \geq \kappa > 0$ (e.g.,
169 $1 + \varepsilon > |\lambda_{\min}(\mathbf{A})|$), $\sigma_{\min}(\mathbf{W}_k) \geq m_k > 0$, and that σ_1, σ_2 are leaky-ReLU with slope $\alpha > 0$ (or
170 ReLU operated with a strict margin $\delta > 0$ along the line segment $\gamma(t) = \mathbf{X} + t\Delta\mathbf{X}$). Then*

$$\|h_\theta(\mathbf{X} + \Delta\mathbf{X}) - h_\theta(\mathbf{X})\|_F \geq c_{\text{lower}} \|\Delta\mathbf{X}\|_F, \quad c_{\text{lower}} = \alpha^2 m_1 m_2 \kappa. \quad (5)$$

171 *A second aggregation step (turning the block into a true two-layer GIN $\sigma_2(\mathbf{M}\sigma_1(\mathbf{M}\mathbf{X}\mathbf{W}_1)\mathbf{W}_2)$)*
172 *replaces c_{lower} by $\alpha^2 m_1 m_2 \kappa^2$; the qualitative content of the lemma is unchanged.*

173 *Proof:* singular-value chain rule on the vectorized Jacobian $\mathbf{J} = \mathbf{D}_2(\mathbf{W}_2^\top \otimes \mathbf{I}_N) \mathbf{D}_1(\mathbf{W}_1^\top \otimes \mathbf{M})$
174 (Appendix B). The lemma encodes that the GIN does not collapse meaningful feature differences, a
175 property the model exploits when fitting noisy labels.

176 6.2 Main Theorem: Noise Fitting Forces Energy Growth

177 Consider two graphs $\mathcal{G}^i, \mathcal{G}^j$ with $N^i = N^j = N$ and similar pooled descriptors $\mathbf{g}_i \approx \mathbf{g}_j$ on the
178 clean distribution (the same- N assumption is for notational simplicity; unequal N in Appendix B).

179 Suppose label noise flips c_j to $c'_j \neq c_i$, forcing the network to output different graph-level labels. Let
 180 $\mathbf{h}_i \in \mathbb{R}^{N \times m'}$ be the post-message-passing signal of $\mathcal{G}^i$, $\mathbf{g}_i = \frac{1}{N} \sum_v \mathbf{h}_{i,v}$, and $\bar{\mathbf{h}}_i = \mathbf{1}_N \mathbf{g}_i^\top$.

181 **Theorem 1** (Graph-level noise fitting forces E^{dir} to grow). *Under the hypotheses of Lemmas 1–2*
 182 *and $N^i = N^j = N$,*

$$\|\mathbf{h}_i - \mathbf{h}_j\|_F \leq \sqrt{N} \|\mathbf{g}_i - \mathbf{g}_j\|_2 + \sqrt{\frac{E^{\text{dir}}(\mathbf{Z}^i)}{\lambda_2(\mathcal{G}^i)}} + \sqrt{\frac{E^{\text{dir}}(\mathbf{Z}^j)}{\lambda_2(\mathcal{G}^j)}}. \quad (6)$$

183 *In particular, if the noisy training objective combined with an L -Lipschitz readout forces a graph-level*
 184 *output gap of at least Δ_{out} , so that $\|\mathbf{h}_i - \mathbf{h}_j\|_F \geq \Theta_{\text{noise}} := L^{-1} \Delta_{\text{out}}$, while the structural similarity*
 185 *gives $\|\mathbf{g}_i - \mathbf{g}_j\|_2 \leq \varepsilon$, then*

$$\sqrt{\frac{E^{\text{dir}}(\mathbf{Z}^i)}{\lambda_2(\mathcal{G}^i)}} + \sqrt{\frac{E^{\text{dir}}(\mathbf{Z}^j)}{\lambda_2(\mathcal{G}^j)}} \geq \Theta_{\text{noise}} - \sqrt{N} \varepsilon. \quad (7)$$

186 **Reading the bound.** (7) lower-bounds the within-graph Dirichlet energies. The intuition is geometric:
 187 for two structurally similar graphs mapped to different labels, $\mathbf{g}_i \approx \mathbf{g}_j$, so the required output gap
 188 Θ_{noise} cannot come from the means alone; by Poincaré, the only remaining route is the within-graph
 189 variance E^{dir}/λ_2 . The unequal- N case (zero-padded matching) is handled in Appendix B. The
 190 bound has no node-classification analog: there labels live on the same nodes whose smoothness E^{dir}
 191 measures, so any link is trivial by construction.

192 **Empirical anchor.** The bound predicts that $\sqrt{E^{\text{dir}}/\lambda_2}$ should grow monotonically with the noise
 193 rate, since higher noise forces the encoder to discriminate more pairs of similar graphs. The PPA
 194 noise-rate sweep in Appendix F confirms this: asymptotic $E_{\text{set}}^{\text{dir}}$ at 10%, 20%, 30%, 40% noise is
 195 monotonically increasing.

196 6.3 Asymmetric Burden on High-Entropy Graphs

197 (7) only constrains the *sum* of $\sqrt{E^{\text{dir}}(\mathbf{Z}^i)/\lambda_2(\mathcal{G}^i)}$ across the two graphs. We next show that the
 198 structural-entropy ranking of Qian et al. [2025] dictates which graph absorbs the burden.

199 For a graph $\mathcal{G}$, the structural Shannon entropy is $H(\mathcal{G}) = -\sum_i \frac{d_i}{\sum_j d_j} \log \frac{d_i}{\sum_j d_j}$. Theorems 1–2
 200 of Qian et al. [2025] prove that more nodes (at fixed sparsity) and denser connectivity (at fixed $|V|$)
 201 both strictly increase H .

202 **Lemma 3** (Spectral counting). *The combinatorial Laplacian $\mathbf{L}^i$ has at most N^i distinct eigenvalues,*
 203 *contained in $[0, \max_{(u,v) \in \mathcal{E}^i} (d_u + d_v)]$, hence in $[0, 2d_{\max}(\mathcal{G}^i)]$ [Chung, 1997]. Equivalently, the*
 204 *normalized spectrum $\Lambda(\Delta^i) \subset [0, 2]$ with $\lambda_{\max}(\Delta^i) = 2$ iff $\mathcal{G}^i$ has a bipartite component. Therefore,*
 205 *when sharpening lifts spectral mass onto eigenvectors ψ_u of $\mathbf{L}$ with $\bar{\lambda}_u \geq \tau$ for some threshold τ , the*
 206 *share allocated to $\mathcal{G}^i$ is bounded above by the count $n^i(\tau) = |\{j : \lambda_j^i \geq \tau\}|$, which is non-decreasing*
 207 *in N^i and in connection density (by Cauchy interlacing under vertex/edge addition).*

208 **Corollary 1** (Asymmetric burden). *If $H(\mathcal{G}^j) \geq H(\mathcal{G}^i)$ via either of the conditions in Qian et al.*
 209 *[2025, Thm. 1–2], then the count of high-frequency eigenvalues of $\mathcal{G}^j$ above any threshold τ is at*
 210 *least that of $\mathcal{G}^i$, so the energy increase forced by Theorem 1 is preferentially absorbed by $\mathcal{G}^j$. This*
 211 *collapses the node-distribution distinctness of $\mathcal{G}^j$ [Qian et al., 2025, Sec. 3.3], the very property*
 212 *pooled graph classifiers rely on.*

213 **Implication.** Theorem 1 says noise fitting must raise energy somewhere; Corollary 1 says it raises
 214 it preferentially in the graphs the classifier most needs to distinguish. Suppressing within-graph
 215 energy growth therefore blocks the noise-fitting path while preserving the classification signal, an
 216 alignment that has no node-classification analog. This also explains why the surge of $E^{\text{dir}}(Z_2)$ under
 217 noise (Section 5) is steeper than under clean overfitting: noise fitting on similar graphs is structurally
 218 constrained to high-frequency, high-entropy directions, while clean overfitting can spread across the
 219 full spectrum.

220 7 Three Methods that Reduce Within-Graph Energy

221 The diagnostic and analytical results above suggest controlling within-graph Dirichlet energy through
 222 any of three mechanistically distinct levers: a regularizer on the energy itself, a loss that down-
 223 weights samples driving the energy up, or an architectural projection that removes the spectral

directions responsible for sharpening. We instantiate one intervention per lever. GCOD is our practical recommendation; CE+W2 is included as a theoretical instantiation.

7.1 Explicit Dirichlet-Energy Regularization ($\mathcal{L}_{DE}$)

The most direct lever is to penalize $E^{\text{dir}}(\mathbf{Z}^i)$ above a class-aware threshold U_{c_i} :

$$\mathcal{L}_{DE}(\mathcal{D}) = \frac{1}{N} \sum_{i=1}^N [\max(0, E^{\text{dir}}(\mathbf{Z}^i) - U_{c_i})]^2, \quad \mathcal{L} = \mathcal{L}_{CE} + \lambda \mathcal{L}_{DE}. \quad (8)$$

We instantiate the threshold in two ways. **Class-specific:** U_c is recomputed every epoch as the mean E^{dir} over a small clean validation subset of class c . This adapts to class-dependent baseline energies (a molecular graph’s natural energy differs from a social graph’s), but requires clean held-out data. **Fixed:** $U_c = U$ globally, a single hyperparameter, no clean validation needed, at the cost of mild over-smoothing on clean labels. Both variants improve robustness; class-specific additionally improves *clean-label* accuracy by acting as a generalization-aware regularizer.

7.2 Graph Centroid Outlier Discounting (GCOD)

GCOD is our practical recommendation. It extends NCOD [Wani et al., 2023] (originally an image-classification loss) to graph classification, with two design changes that are necessary in the graph-classification regime: a third KL term that disentangles clean from noisy graphs via per-sample trainable parameters $\mathbf{u} \in [0, 1]^n$, and feedback through the current training accuracy a_{train} that prevents premature outlier discounting. For batch B with predictions $f_{\theta}(\mathbf{Z}_B)$, hard predictions $\hat{\mathbf{y}}_B$, soft labels $\tilde{\mathbf{y}}_B$ (centroid-based, see Appendix C.6), and observed labels $\mathbf{y}_B$:

$$\mathcal{L}_1 = \mathcal{L}_{CE}(f_{\theta}(\mathbf{Z}_B) + a_{\text{train}} \text{diag}_{\text{vec}}(\mathbf{u}_B) \mathbf{y}_B, \tilde{\mathbf{y}}_B), \quad (9)$$

$$\mathcal{L}_2 = \frac{1}{|C|} \|\hat{\mathbf{y}}_B + \text{diag}_{\text{vec}}(\mathbf{u}_B) \mathbf{y}_B - \mathbf{y}_B\|^2, \quad (10)$$

$$\mathcal{L}_3 = (1 - a_{\text{train}}) D_{KL}\{\mathcal{L}, \sigma(-\log \mathbf{u}_B)\}, \quad (11)$$

with $\mathcal{L} = \log \sigma(\text{diag}_{\text{mat}}(f_{\theta}(\mathbf{Z}_B) \mathbf{y}_B^{\top}))$. Updates are split: $\theta \leftarrow \theta - \alpha \nabla_{\theta}(\mathcal{L}_1 + \mathcal{L}_3)$, $\mathbf{u} \leftarrow \mathbf{u} - \beta \nabla_{\mathbf{u}} \mathcal{L}_2$. GCOD is not a theoretically derived energy minimizer. Empirically, however, it reduces $E_{\text{set}}^{\text{dir}}$ during training (Fig. 2) without explicitly penalizing it: by down-weighting graphs flagged as noisy via large u_i , it removes the gradient pressure that would otherwise push the network toward high-frequency directions, exactly as Theorem 1 predicts for any sample-level mechanism that suppresses noise-driven gradients.

7.3 Spectral Constraint on Weight Matrices (CE+W2)

A third lever, included as a theoretical instantiation, acts directly on the weights. Giovanni et al. [2023] show that for linear graph convolutions with symmetric $\mathbf{W}$, the gradient flow of E^{dir} decomposes along eigenvectors of $\mathbf{W}$: positive eigenvalues smooth, negative ones sharpen. Practical $\mathbf{W}_l^2$ is not symmetric, so we split $\mathbf{W}_l^2 = \mathbf{S}_l + \mathbf{N}_l$ and project only $\mathbf{S}_l$ to its positive eigenspace, leaving the skew $\mathbf{N}_l$ (purely imaginary spectrum, rotational) untouched. This is a heuristic extension of GRAFF — the energy of $\mathbf{MZW}$ depends on $\mathbf{W}^{\top} \mathbf{W}$, not just $\mathbf{S}_l$ — but the resulting CE+W2 needs neither clean validation nor extra hyperparameters and shows the energy mechanism can be controlled architecturally. Implementation details are in Appendix C.4; the per-epoch eigendecomposition is costly and can destabilize convergence, so we recommend $\mathcal{L}_{DE}$ or GCOD for practical use.

8 Experiments

Setup. We evaluate on ogbg-ppa [Szklarczyk et al., 2018a] (full and a controlled 30%/6-class subset), MUTAG, PROTEINS, ENZYMES, IMDB-BINARY, REDDIT-MULTI-12K, MSRC-21, MNIST-graph, and Cora (node classification, transfer test). Backbones are GIN [Xu et al., 2019] and GCN [Kipf and Welling, 2017]; full hyperparameters are in Appendix C. We compare against three baselines: standard cross-entropy (CE); SOP [Liu et al., 2022], a strong noise-robust loss from

Table 1: Test accuracy on the PPA-30%-6 subset (5-layer GIN, mean $\pm$ std over 4 runs). **Best** is highlighted in red and bold, second-best in blue. “CE clean only” shows accuracy on the clean subset of the noisy data, as a reference upper bound. Class-specific $\mathcal{L}_{DE}^*$ requires a clean validation set.

Noise	Method	Test Acc.		Train Acc.	
		Best	Last	Best	Last
0%	CE	96.25 $\pm$ 0.05	91.25	100.0	99.33
	GCOD	96.65 $\pm$ 0.52	93.25	99.23	99.03
	CE+W2	96.50 $\pm$ 1.12	86.50	99.76	99.29
	$\mathcal{L}_{DE}$ (Fixed)	96.58 $\pm$ 0.27	85.70	99.26	98.29
	$\mathcal{L}_{DE}^*$ (Class-spec.)	96.96 $\pm$ 0.04	88.67	99.41	98.67
20%	CE (clean only)	96.15 $\pm$ 0.09	90.66	84.50	84.07
	CE	88.66 $\pm$ 0.16	62.58	94.98	93.45
	SOP	91.01 $\pm$ 0.44	85.50	79.17	77.64
	GCOD	93.91 $\pm$ 0.26	92.58	80.11	79.52
	CE+W2	89.83 $\pm$ 1.23	76.66	84.90	83.68
	$\mathcal{L}_{DE}$ (Fixed)	89.69 $\pm$ 0.57	80.25	78.07	70.14
	$\mathcal{L}_{DE}^*$	92.34 $\pm$ 0.41	80.92	84.55	83.73
40%	CE (clean only)	95.08 $\pm$ 0.13	80.44	68.15	67.05
	CE	82.08 $\pm$ 0.22	51.83	82.47	78.88
	SOP	82.33 $\pm$ 1.32	65.41	58.90	57.68
	GCOD	93.88 $\pm$ 0.04	91.08	65.09	64.31
	CE+W2	88.25 $\pm$ 1.13	56.33	68.78	65.88
	$\mathcal{L}_{DE}$ (Fixed)	88.58 $\pm$ 0.75	77.12	61.08	54.53
	$\mathcal{L}_{DE}^*$	88.55 $\pm$ 0.11	71.83	68.98	67.80

image classification; and OMG [Yin et al., 2023], the leading graph-classification noise method. All numbers are mean $\pm$ std over 4 runs.

Why a controlled PPA subset? On the full 37-class PPA, all methods are under-parameterized for the task, and CE itself is robust (Section 4); a noise-robustness method cannot help where the baseline does not fail. We use a 30%/6-class PPA subset as a *controlled* testbed where over-parameterization triggers noise fitting, and we additionally report the full-PPA result to show our methods do no harm there. Generality claims are grounded in the seven additional benchmarks.

8.1 Main Result: Controlled PPA Subset

Table 1 reports the full ablation on the controlled PPA subset. GCOD is best at every noise level and, crucially, has the smallest gap between best and final test accuracy, indicating that it does not overfit late in training. CE+W2 and $\mathcal{L}_{DE}$ both improve robustness over CE, confirming the convergent-evidence prediction of Section 6. Under clean labels, $\mathcal{L}_{DE}^*$ even outperforms CE, a direct effect of the class-aware energy threshold.

8.2 Generalization Across Datasets and Noise Types

Table 2 reports GCOD across seven benchmarks beyond PPA. GCOD consistently improves over CE under both 20% symmetric and 40% asymmetric noise, and matches CE on REDDIT where the baseline is already saturated by class imbalance. The empirical trend that gains grow with the noise rate (Appendix F) is consistent with Theorem 1: the forced energy lower bound $\Theta_{\text{noise}} - \sqrt{N}\varepsilon$ scales with the noise level, so the room a method has to help also scales.

Table 2: GCOD matches or beats CE across diverse datasets. Mean over 4 runs; $\pm$ std reported where available. *Sym*₂₀: 20% symmetric noise. *Asym*₄₀: 40% asymmetric noise. The *Asym*₄₀ GCOD column reports the std deviations from [Wani et al., 2023]’s rebuttal protocol.

Dataset	0% CE	Sym ₂₀ CE	Sym ₂₀ GCOD	Asym ₄₀ CE	Asym ₄₀ GCOD
PROTEINS	81.16	76.23	79.38	72.90	76.19 $\pm$ 1.22
MNIST	72.69	66.30	71.26	71.15	72.61 $\pm$ 0.43
ENZYMES	73.33	64.16	68.69	65.80	69.81 $\pm$ 0.98
IMDB-B	76.50	75.00	75.40	68.18	72.89 $\pm$ 0.91
MUTAG	94.73	89.47	90.01	91.16	93.19 $\pm$ 0.38
REDDIT	48.15	45.05	44.98	48.01	47.94 $\pm$ 2.87
MSRC-21	96.69	90.26	94.69	94.39	95.57 $\pm$ 0.32

Table 3: OMG protocol (Yin et al., 2023). Percentage gain over CE.

Dataset	OMG	GCOD
MUTAG	0.061	0.062
IMDB-B	0.047	0.049
PROTEINS	0.039	0.041

Table 4: Relative training time (CE-GIN = 1.00).

Method	Runtime
CE-GIN	1.00
SOP	1.048
GCOD	1.029
CE+W2	1.33
OMG [Yin et al., 2023]	$\approx$ 3.00

282 **Full ogbg-ppa.** Under the full 37-class benchmark, all methods including CE behave robustly to
283 noise (Section 4); we therefore use it only as a do-no-harm check. Detailed results are in Appendix H.

284 **Comparison with OMG.** Under the OMG [Yin et al., 2023] protocol (Table 3), GCOD matches
285 OMG’s accuracy (within 0.1–0.2 pp) while requiring $\sim 70\times$ less training overhead (Table 4): OMG
286 adds $\approx 200\%$ runtime, GCOD only $\approx 2.9\%$.

287 **Cross-setting transfer to node classification (Cora).** Applied unchanged to node classification
288 on Cora under 20% symmetric noise (full table in Appendix I, against 13 node-specific baselines
289 including NRGNN, RTGNN, GNN-Cleaner, ERASE, CR-GNN), our methods rank in the top
290 half on both GAT and GCN backbones: CE+W2 reaches 0.8157 (rank #2) on GAT and 0.7247
291 (rank #5) on GCN; GCOD reaches 0.7910 (rank #4) on GAT and 0.7403 (rank #4) on GCN. An
292 energy-based regularizer that uses neither neighbor labels nor homophily yet outperforms specialized
293 node-classification methods is direct evidence that the perspective captures a general property of
294 GNN training under label noise.

295 **Convergent evidence.** The three methods are mechanistically distinct (architectural, regularization,
296 sample re-weighting), yet all reduce $E_{\text{set}}^{\text{dir}}$ during training (Fig. 2, Appendix F) and improve robustness
297 over CE — three independent levers acting on the same target, as Theorem 1 and Corollary 1 predict.

298 9 Conclusion

299 Within-graph Dirichlet energy of learned representations is an empirical signal of graph-level noise
300 memorization, supported by a Poincaré–bi-Lipschitz lower bound and an entropy-based asymmetry
301 argument specific to graph classification. The diagnostic translates into three interventions ($\mathcal{L}_{DE}$,
302 GCOD, CE+W2) that all reduce within-graph energy and improve robustness across eight benchmarks
303 while preserving clean-label performance; GCOD adds $\approx 3\%$ training overhead and matches OMG at
304 $\sim 70\times$ lower compute. **Limitations.** The bound is a sufficient condition, not a closed-form character-
305 ization; we use synthetic noise protocols since real-world noisy graph-classification benchmarks are
306 scarce; CE+W2 is a heuristic extension of GRAFF. **Broader impact.** Robust graph learning matters
307 in chemistry and biology, where annotation errors are systematic, and GCOD avoids the need for
308 clean validation data.

References

- Görkem Algan and Ilkay Ulusoy. Image classification with deep learning in the presence of noisy labels: A survey. *Knowledge-Based Systems*, 215:106771, 2021. ISSN 0950-7051. doi: <https://doi.org/10.1016/j.knosys.2021.106771>. URL <https://www.sciencedirect.com/science/article/pii/S0950705121000344>.
- Eric Arazo, Diego Ortego, Paul Albert, Noel O'Connor, and Kevin McGuinness. Unsupervised label noise modeling and loss correction. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 312–321. PMLR, 09–15 Jun 2019. URL <https://proceedings.mlr.press/v97/arazo19a.html>.
- Devansh Arpit, Stanisław Jastrzębski, Nicolas Ballas, David Krueger, Emmanuel Bengio, Maxinder S. Kanwal, Tegan Maharaj, Asja Fischer, Aaron Courville, Yoshua Bengio, and Simon Lacoste-Julien. A closer look at memorization in deep networks. In Doina Precup and Yee Whye Teh, editors, *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 233–242. PMLR, 06–11 Aug 2017. URL <https://proceedings.mlr.press/v70/arpit17a.html>.
- Deyu Bo, Xiao Wang, Chuan Shi, and Huawei Shen. Beyond low-frequency information in graph convolutional networks. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(5): 3950–3957, May 2021. doi: 10.1609/aaai.v35i5.16514. URL <https://ojs.aaai.org/index.php/AAAI/article/view/16514>.
- Karsten M. Borgwardt, Cheng Soon Ong, Stefan Schönauer, S. V. N. Vishwanathan, Alex J. Smola, and Hans-Peter Kriegel. Protein function prediction via graph kernels. *Bioinformatics*, 21(suppl_1): i47–i56, 06 2005. ISSN 1367-4803. doi: 10.1093/bioinformatics/bti1007. URL <https://doi.org/10.1093/bioinformatics/bti1007>.
- Chen Cai and Yusu Wang. A note on over-smoothing for graph neural networks, 2020. URL <https://arxiv.org/abs/2006.13318>.
- Jialin Chen, Yuelin Wang, Cristian Bodnar, Rex Ying, Pietro Lio, and Yu Guang Wang. Dirichlet energy enhancement of graph neural networks by framelet augmentation, 2023. URL <https://arxiv.org/abs/2311.05767>.
- Ching-Yao Chuang and Stefanie Jegelka. Tree mover's distance: Bridging graph metrics and stability of graph neural networks. *proceedings neurips*, 35:2944–2957, 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/file/139ae969f49abd9a113981c1f7fce5ce-Paper-Conference.pdf.
- Fan R. K. Chung. *Spectral Graph Theory*. Number 92 in CBMS Regional Conference Series in Mathematics. American Mathematical Society, 1997.
- Enyan Dai, Charu Aggarwal, and Suhang Wang. Nrgnn: Learning a label noise resistant graph neural network on sparsely and noisily labeled graphs. In *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, KDD '21, page 227–236, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450383325. doi: 10.1145/3447548.3467364. URL <https://doi.org/10.1145/3447548.3467364>.
- Enyan Dai, Wei Jin, Hui Liu, and Suhang Wang. Towards robust graph neural networks for noisy graphs with sparse labels. In *Proceedings of the Fifteenth ACM International Conference on Web Search and Data Mining*, WSDM '22, page 181–191, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450391320. doi: 10.1145/3488560.3498408. URL <https://doi.org/10.1145/3488560.3498408>.
- Yair Davidson and Nadav Dym. On the holder stability of multiset and graph neural networks. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=P7KIGdgW8S>.

357 Xuefeng Du, Tian Bian, Yu Rong, Bo Han, Tongliang Liu, Tingyang Xu, Wenbing Huang, Yixuan
358 Li, and Junzhou Huang. Noise-robust graph learning by estimating and leveraging pairwise
359 interactions. *Transactions on Machine Learning Research*, 2023. ISSN 2835-8856. URL <https://openreview.net/forum?id=r7imkFEAQb>.
360

361 James Fox and Sivasankaran Rajamanickam. How robust are graph neural networks to structural
362 noise? *ArXiv*, abs/1912.10206, 2019. URL [https://api.semanticscholar.org/CorpusID:](https://api.semanticscholar.org/CorpusID:209444207)
363 209444207.

364 Fernando Gama, Joan Bruna, and Alejandro Ribeiro. Stability properties of graph neural net-
365 works. *IEEE Transactions on Signal Processing*, 68:5680–5695, 2019. URL [https://api.](https://api.semanticscholar.org/CorpusID:152282530)
366 [semanticscholar.org/CorpusID:152282530](https://api.semanticscholar.org/CorpusID:152282530).

367 Aritra Ghosh, Himanshu Kumar, and P Shanti Sastry. Robust loss functions under label noise for
368 deep neural networks. In *Proceedings of the AAAI conference on artificial intelligence*, volume 31,
369 2017.

370 Francesco Di Giovanni, James Rowbottom, Benjamin Paul Chamberlain, Thomas Markovich, and
371 Michael M. Bronstein. Understanding convolution on graphs via energies. *Transactions on*
372 *Machine Learning Research*, 2023. ISSN 2835-8856. URL [https://openreview.net/forum?](https://openreview.net/forum?id=v5ew3FPTgb)
373 [id=v5ew3FPTgb](https://openreview.net/forum?id=v5ew3FPTgb).

374 Xian-Jin Gui, Wei Wang, and Zhang-Hao Tian. Towards understanding deep learning from noisy
375 labels with small-loss criterion. In Zhi-Hua Zhou, editor, *Proceedings of the Thirtieth Interna-*
376 *tional Joint Conference on Artificial Intelligence, IJCAI-21*, pages 2469–2475. International Joint
377 Conferences on Artificial Intelligence Organization, 8 2021. doi: 10.24963/ijcai.2021/340. URL
378 <https://doi.org/10.24963/ijcai.2021/340>. Main Track.

379 Bo Han, Quanming Yao, Xingrui Yu, Gang Niu, Miao Xu, Weihua Hu, Ivor Tsang, and Masashi
380 Sugiyama. Co-teaching: Robust training of deep neural networks with extremely noisy labels.
381 *Advances in neural information processing systems*, 31, 2018.

382 Roger A. Horn and Charles R. Johnson. *Matrix Analysis*. Cambridge University Press, 2 edition,
383 2012.

384 Simona Juvina, Ana-Antonia Neacșu, Jean-Christophe Pesquet, Burileanu Corneliu, Jérôme Rony,
385 and Ismail Ben Ayed. Training graph neural networks subject to a tight lipschitz constraint.
386 *Transactions on Machine Learning Research Journal*, 2024.

387 Zhao Kang, Haiqi Pan, Steven C. H. Hoi, and Zenglin Xu. Robust graph learning from
388 noisy data. *IEEE Transactions on Cybernetics*, 50:1833–1843, 2018. URL [https://api.](https://api.semanticscholar.org/CorpusID:56208557)
389 [semanticscholar.org/CorpusID:56208557](https://api.semanticscholar.org/CorpusID:56208557).

390 Jihye Kim, Aristide Baratin, Yan Zhang, and Simon Lacoste-Julien. Crosssplit: Mitigating label
391 noise memorization through data splitting. In *International Conference on Machine Learning*,
392 pages 16377–16392. PMLR, 2023.

393 Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional networks.
394 In *International Conference on Learning Representations*, 2017. URL [https://openreview.](https://openreview.net/forum?id=SJU4ayYgl)
395 [net/forum?id=SJU4ayYgl](https://openreview.net/forum?id=SJU4ayYgl).

396 Nils M. Kriege and Petra Mutzel. Subgraph matching kernels for attributed graphs. In *Interna-*
397 *tional Conference on Machine Learning*, 2012. URL [https://api.semanticscholar.org/](https://api.semanticscholar.org/CorpusID:8790929)
398 [CorpusID:8790929](https://api.semanticscholar.org/CorpusID:8790929).

399 Junnan Li, Richard Socher, and Steven C. H. Hoi. Dividemix: Learning with noisy labels as semi-
400 supervised learning. In *ICLR*, 2020. URL <https://openreview.net/forum?id=HJgExaVtwr>.

401 Xianxian Li, Qiyu Li, De Li, Haodong Qian, and Jinyan Wang. Contrastive learning of graphs
402 under label noise. *Neural Networks*, 172:106113, 2024. ISSN 0893-6080. doi: [https://doi.org/](https://doi.org/10.1016/j.neunet.2024.106113)
403 [10.1016/j.neunet.2024.106113](https://doi.org/10.1016/j.neunet.2024.106113). URL [https://www.sciencedirect.com/science/article/](https://www.sciencedirect.com/science/article/pii/S0893608024000273)
404 [pii/S0893608024000273](https://www.sciencedirect.com/science/article/pii/S0893608024000273).

405 Soon Hoe Lim, N. Benjamin Erichson, Francisco Utrera, Winnie Xu, and Michael W. Mahoney.
 406 Noisy feature mixup. In *International Conference on Learning Representations*, 2022. URL
 407 <https://openreview.net/forum?id=vJb4I2ANmy>.

408 Sheng Liu, Zhihui Zhu, Qing Qu, and Chong You. Robust training under label noise by over-
 409 parameterization. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu,
 410 and Sivan Sabato, editors, *Proceedings of the 39th International Conference on Machine Learning*,
 411 volume 162 of *Proceedings of Machine Learning Research*, pages 14153–14172. PMLR, 17–23
 412 Jul 2022. URL <https://proceedings.mlr.press/v162/liu22w.html>.

413 Tongliang Liu and Dacheng Tao. Classification with noisy labels by importance reweighting. *IEEE*
 414 *Transactions on pattern analysis and machine intelligence*, 38(3):447–461, 2015.

415 Marion Neumann, Roman Garnett, Christian Bauckhage, and Kristian Kersting. Propagation ker-
 416 nels: Efficient graph kernels from propagated information. *Machine Learning*, 102(2):209–245,
 417 2016. ISSN 1573-0565. doi: 10.1007/s10994-015-5517-9. URL [https://doi.org/10.1007/](https://doi.org/10.1007/s10994-015-5517-9)
 418 [s10994-015-5517-9](https://doi.org/10.1007/s10994-015-5517-9).

419 Hoang NT and Takanori Maehara. Revisiting graph neural networks: All we have is low-pass
 420 filters. *ArXiv*, abs/1905.09550, 2019. URL [https://api.semanticscholar.org/CorpusID:](https://api.semanticscholar.org/CorpusID:162183860)
 421 [162183860](https://api.semanticscholar.org/CorpusID:162183860).

422 Hoang NT, Jun Jin Choong, and Tsuyoshi Murata. Learning graph neural networks with noisy labels,
 423 2019. URL <https://openreview.net/forum?id=r1x0mNmxuN>.

424 Kenta Oono and Taiji Suzuki. Graph neural networks exponentially lose expressive power for
 425 node classification. In *International Conference on Learning Representations*, 2020. URL
 426 <https://openreview.net/forum?id=S1ld02EFPr>.

427 Giorgio Patrini, Alessandro Rozza, Aditya Krishna Menon, Richard Nock, and Lizhen Qu. Making
 428 deep neural networks robust to label noise: A loss correction approach. In *Proceedings of the*
 429 *IEEE conference on computer vision and pattern recognition*, pages 1944–1952, 2017.

430 Geoff Pleiss, Tianyi Zhang, Ethan Elenberg, and Kilian Q Weinberger. Identifying mislabeled data
 431 using the area under the margin ranking. *Advances in Neural Information Processing Systems*, 33:
 432 17044–17056, 2020.

433 Feifei Qian, Lu Bai, Lixin Cui, Ming Li, Hangyuan Du, Yue Wang, and Edwin Hancock. Exploring
 434 the over-smoothing problem of graph neural networks for graph classification: An entropy-based
 435 viewpoint. In *Proceedings of the Thirty-Fourth International Joint Conference on Artificial*
 436 *Intelligence (IJCAI)*, pages 3235–3243, 2025.

437 Siyi Qian, Haochao Ying, Renjun Hu, Jingbo Zhou, Jintai Chen, Danny Z. Chen, and Jian Wu.
 438 Robust training of graph neural networks via noise governance. In *Proceedings of the Sixteenth*
 439 *ACM International Conference on Web Search and Data Mining, WSDM ’23*, page 607–615,
 440 New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9781450394079. doi:
 441 [10.1145/3539597.3570369](https://doi.org/10.1145/3539597.3570369). URL <https://doi.org/10.1145/3539597.3570369>.

442 Scott E. Reed, Honglak Lee, Dragomir Anguelov, Christian Szegedy, Dumitru Erhan, and Andrew
 443 Rabinovich. Training deep neural networks on noisy labels with bootstrapping. In *ICLR 2015*,
 444 2015. URL <http://arxiv.org/abs/1412.6596>.

445 T. Konstantin Rusch, Michael M. Bronstein, and Siddhartha Mishra. A survey on oversmoothing in
 446 graph neural networks. *ArXiv*, abs/2303.10993, 2023. URL [https://api.semanticscholar.](https://api.semanticscholar.org/CorpusID:257632346)
 447 [org/CorpusID:257632346](https://api.semanticscholar.org/CorpusID:257632346).

448 Damian Szklarczyk, Annika Gable, David Lyon, Alexander Junge, Stefan Wyder, Jaime Huerta-
 449 Cepas, Milan Simonovic, Nadezhda Doncheva, John Morris, Peer Bork, Lars Jensen, and Christian
 450 von Mering. String v11: protein-protein association networks with increased coverage, supporting
 451 functional discovery in genome-wide experimental datasets. *Nucleic acids research*, 47, 11 2018a.
 452 doi: 10.1093/nar/gky1131.

453 Damian Szklarczyk, Annika L Gable, David Lyon, Alexander Junge, Stefan Wyder, Jaime Huerta-
454 Cepas, Milan Simonovic, Nadezhda T Doncheva, John H Morris, Peer Bork, Lars J Jensen,
455 and Christian von Mering. STRING v11: protein–protein association networks with increased
456 coverage, supporting functional discovery in genome-wide experimental datasets. *Nucleic Acids*
457 *Research*, 47(D1):D607–D613, 11 2018b. ISSN 0305-1048. doi: 10.1093/nar/gky1131. URL
458 <https://doi.org/10.1093/nar/gky1131>.

459 Nikil Wale and George Karypis. Comparison of descriptor spaces for chemical compound retrieval
460 and classification. In *Sixth International Conference on Data Mining (ICDM'06)*, pages 678–689,
461 2006. doi: 10.1109/ICDM.2006.39.

462 Farooq Ahmad Wani, Maria Sofia Bucarelli, and Fabrizio Silvestri. Learning with noisy labels
463 through learnable weighting and centroid similarity. *2024 International Joint Conference on Neural*
464 *Networks (IJCNN)*, pages 1–9, 2023. URL [https://api.semanticscholar.org/CorpusID:](https://api.semanticscholar.org/CorpusID:257557810)
465 257557810.

466 Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural
467 networks? In *7th International Conference on Learning Representations, ICLR 2019, New Orleans,*
468 *LA, USA, May 6-9, 2019*. OpenReview.net, 2019. URL [https://openreview.net/forum?id=](https://openreview.net/forum?id=ryGs6iA5Km)
469 ryGs6iA5Km.

470 Shilong Xu, Jipeng Guo, Jiahao Long, Mingliang Cui, and Youqing Wang. High-low frequency
471 filtering aware graph neural networks with representation decoupling learning. In *2024 China*
472 *Automation Congress (CAC)*, pages 4005–4010, 2024. doi: 10.1109/CAC63892.2024.10864603.

473 Yujun Yan, Milad Hashemi, Kevin Swersky, Yaoqing Yang, and Danai Koutra. Two sides of
474 the same coin: Heterophily and oversmoothing in graph convolutional neural networks. In
475 *2022 IEEE International Conference on Data Mining (ICDM)*, pages 1287–1292, 2022. doi:
476 10.1109/ICDM54844.2022.00169.

477 Pinar Yanardag and S.V.N. Vishwanathan. Deep graph kernels. In *Proceedings of the 21th*
478 *ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '15,
479 page 1365–1374, New York, NY, USA, 2015a. Association for Computing Machinery. ISBN
480 9781450336642. doi: 10.1145/2783258.2783417. URL [https://doi.org/10.1145/2783258.](https://doi.org/10.1145/2783258.2783417)
481 2783417.

482 Pinar Yanardag and S.V.N. Vishwanathan. Deep graph kernels. In *Proceedings of the 21th*
483 *ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '15,
484 page 1365–1374, New York, NY, USA, 2015b. Association for Computing Machinery. ISBN
485 9781450336642. doi: 10.1145/2783258.2783417. URL [https://doi.org/10.1145/2783258.](https://doi.org/10.1145/2783258.2783417)
486 2783417.

487 Nan Yin, Li Shen, Mengzhu Wang, Xiao Luo, Zhigang Luo, and Dacheng Tao. Omg: Towards
488 effective graph classification against label noise. *IEEE Transactions on Knowledge and Data*
489 *Engineering*, 35(12):12873–12886, 2023. doi: 10.1109/TKDE.2023.3271677.

490 Jingyang Yuan, Xiao Luo, Yifang Qin, Yusheng Zhao, Wei Ju, and Ming Zhang. Learning on graphs
491 under label noise. In *ICASSP 2023 - 2023 IEEE International Conference on Acoustics, Speech*
492 *and Signal Processing (ICASSP)*, pages 1–5, 2023a. doi: 10.1109/ICASSP49357.2023.10096088.

493 Jinliang Yuan, Hualei Yu, Meng Cao, Jianqing Song, Junyuan Xie, and Chongjun Wang. Self-
494 supervised robust graph neural networks against noisy graphs and noisy labels. *Applied Intelligence*,
495 53(21):25154–25170, aug 2023b. ISSN 0924-669X. doi: 10.1007/s10489-023-04836-6. URL
496 <https://doi.org/10.1007/s10489-023-04836-6>.

497 Chiyuan Zhang, Samy Bengio, Moritz Hardt, Benjamin Recht, and Oriol Vinyals. Understanding
498 deep learning (still) requires rethinking generalization. *Commun. ACM*, 64(3):107–115, February
499 2021. ISSN 0001-0782. doi: 10.1145/3446776. URL <https://doi.org/10.1145/3446776>.

500 Hongyi Zhang, Moustapha Ciss'e, Yann N. Dauphin, and David Lopez-Paz. mixup: Beyond empirical
501 risk minimization. *CoRR*, abs/1710.09412, 2017.

- 502 Dengyong Zhou and Bernhard Schölkopf. Regularization on discrete spaces. In Walter G. Kropatsch,
503 Robert Sablatnig, and Allan Hanbury, editors, *Pattern Recognition*, pages 361–368, Berlin, Heidel-
504 berg, 2005. Springer Berlin Heidelberg. ISBN 978-3-540-31942-9.
- 505 Kaixiong Zhou, Xiao Huang, Daochen Zha, Rui Chen, Li Li, Soo-Hyun Choi, and Xia Hu. Dirichlet
506 energy constrained learning for deep graph neural networks. *Advances in neural information*
507 *processing systems*, 34:21834–21846, 2021.

A Additional Related Work

A.1 Learning under Label Noise

Beyond robust losses, additional families include sample relabelling [Arazo et al., 2019, Reed et al., 2015], two-network co-teaching and divide-and-mix approaches [Han et al., 2018, Li et al., 2020, Kim et al., 2023], and regularization through MixUp [Zhang et al., 2017]. Reweighting methods adapt sample importance via training dynamics or held-out probes [Liu and Tao, 2015, Pleiss et al., 2020]. Most of these were developed for image classification; their direct transfer to graph classification is non-obvious because graphs lack the strong feature regularities of images.

A.2 Graph Learning under Noise

Most graph-related noise work targets node classification, exploiting either homophily (RS-GNN [Dai et al., 2022], NRGNN [Dai et al., 2021], RTGNN [Yuan et al., 2023a], pairwise interactions [Du et al., 2023]) or self-supervision (pseudo-label denoising [Yuan et al., 2023b, Qian et al., 2023]). Edge/feature noise has also been studied, e.g. structural noise [Fox and Rajamanickam, 2019, Dai et al., 2022]. These methods typically assume connected nodes share labels, which is incompatible with our graph-classification setting where labels are graph-level. The seminal work of NT et al. [2019] treats graph classification under label noise via a surrogate loss; OMG [Yin et al., 2023] combines contrastive learning, MixUp, and curriculum filtering. Both are practical improvements but provide no spectral or energy-based explanation of *why* GNNs fit graph-level noise.

A.3 Dirichlet Energy and Spectral Properties of GNNs

NT and Maehara [2019] and Rusch et al. [2023] analyze GNNs as low-pass filters, with self-loops further shrinking eigenvalues. Cai and Wang [2020], Oono and Suzuki [2020] prove that GNN Dirichlet energy decays exponentially when the product of the largest singular value of the weight matrix and the largest Laplacian eigenvalue is below 1. Giovanni et al. [2023] disentangle smoothing/sharpening via positive/negative weight-eigenvalues. Zhou et al. [2021] (NeurIPS) uses energy regularization to prevent over-smoothing on *node* classification. Our setting is the opposite regime: too *much* energy under graph-level label noise.

Qian et al. [2025] introduce a structural-entropy viewpoint on over-smoothing for graph classification, showing that high-entropy regions are most damaged by over-smoothing. We invert and combine this with Proposition 1 to get Corollary 1: the graphs that suffer most from over-*sharpening* (the noise-fitting regime) are also the high-entropy graphs that carry the most discriminative signal. The two papers therefore identify symmetric ends of the same spectral failure.

A.4 Lipschitz/Stability Analyses

Chuang and Jegelka [2022] bound GIN outputs via the Tree Mover’s Distance, and Davidson and Dym [2025] extend uniform-Lipschitz analyses to MPNNs through Hölder continuity of aggregation/combination/readout. Juvina et al. [2024] derive optimal Lipschitz constants $\vartheta = \phi(\lambda_K)$ for graph convolutions under non-negative weights and ReLU. Gama et al. [2019] show output stability under graph-shift-operator perturbations. These results characterize stability w.r.t. inputs; we instead study stability against *label* noise, mediated by the within-graph energy of learned representations.

B Proofs and Extended Theory

B.1 Proof of Proposition 1

Let $\Lambda^i = \{\lambda_u^i\}_{u=1}^{N^i}$ be the spectrum of the per-graph combinatorial Laplacian $\mathbf{L}^i = \mathbf{D}^i - \mathbf{A}^i$. Define the block-diagonal Laplacian

$$\bar{\mathbf{L}} = \text{blkdiag}(\mathbf{L}^1, \dots, \mathbf{L}^n), \quad \det(\bar{\mathbf{L}} - \lambda \mathbf{I}_{N_{\text{tot}}}) = \prod_{i=1}^n \det(\mathbf{L}^i - \lambda \mathbf{I}_{N^i}), \quad (12)$$

so $\bar{\Lambda} = \bigcup_i \Lambda^i$. Let ψ_j^i be the j -th eigenvector of $\mathbf{L}^i$ at eigenvalue λ_j^i . The vector

$$\bar{\psi}_u = (0, \dots, 0, \psi_j^i, 0, \dots, 0)^\top \in \mathbb{R}^{N_{\text{tot}}} \quad (13)$$

(zero on every block other than i) satisfies $\bar{\mathbf{L}}\bar{\psi}_u = \lambda_j^i \bar{\psi}_u$. With $\bar{\mathbf{Z}} = [\mathbf{Z}^1 \parallel \dots \parallel \mathbf{Z}^n]$ (vertical concatenation), the combinatorial Dirichlet energy of the block-diagonal graph satisfies

$$E^{\text{dir}}(\bar{\mathbf{Z}}) = \text{tr}(\bar{\mathbf{Z}}^\top \bar{\mathbf{L}} \bar{\mathbf{Z}}) = \sum_{r=1}^{m'} \sum_{u=1}^{N_{\text{tot}}} \bar{\lambda}_u (\bar{\psi}_u^\top \bar{\mathbf{Z}}_{:,r})^2 = \sum_i \sum_{r,j} \lambda_j^i (\psi_j^{i\top} \mathbf{Z}_{:,r}^i)^2 = \sum_i E^{\text{dir}}(\mathbf{Z}^i), \quad (14)$$

where the second equality uses the spectral expansion in the orthonormal basis $\{\bar{\psi}_u\}$ and the third uses the block-supported structure of each $\bar{\psi}_u$. Hence $E_{\text{set}}^{\text{dir}}(\mathcal{D}) = \frac{1}{|\mathcal{D}|} \sum_i E^{\text{dir}}(\mathbf{Z}^i) = \frac{1}{|\mathcal{D}|} E^{\text{dir}}(\bar{\mathbf{Z}})$. $\square$

B.2 Proof of Lemma 1 (Graph Poincaré Inequality)

The combinatorial Laplacian $\mathbf{L} = \mathbf{D} - \mathbf{A}$ is real symmetric and positive semi-definite. Because $\mathcal{G}$ is connected, $\ker(\mathbf{L}) = \text{span}(\mathbf{1})$, and the eigenvalues are $0 = \lambda_1 < \lambda_2 \leq \dots \leq \lambda_N$. By the Courant–Fischer min–max characterization,

$$\lambda_2(\mathcal{G}) = \min_{\substack{\mathbf{x} \perp \mathbf{1} \\ \mathbf{x} \neq \mathbf{0}}} \frac{\mathbf{x}^\top \mathbf{L} \mathbf{x}}{\mathbf{x}^\top \mathbf{x}}. \quad (15)$$

Fix any feature dimension r and let $\mathbf{z} = \mathbf{Z}_{:,r} \in \mathbb{R}^N$ with $\bar{z} = \frac{1}{N} \sum_v z_v$. The centered vector $\mathbf{z} - \bar{z}\mathbf{1}$ is orthogonal to $\mathbf{1}$, so

$$\lambda_2(\mathcal{G}) \|\mathbf{z} - \bar{z}\mathbf{1}\|_2^2 \leq (\mathbf{z} - \bar{z}\mathbf{1})^\top \mathbf{L} (\mathbf{z} - \bar{z}\mathbf{1}) = \mathbf{z}^\top \mathbf{L} \mathbf{z} = \sum_{(u,v) \in \mathcal{E}} (z_u - z_v)^2, \quad (16)$$

using $\mathbf{L}\mathbf{1} = \mathbf{0}$. The right-hand side is the per-feature-dimension contribution to $E^{\text{dir}}(\mathbf{Z})$ in (1). Summing over $r = 1, \dots, m'$ yields

$$\sum_{v=1}^N \|\mathbf{Z}_v - \bar{\mathbf{z}}\|_2^2 = \sum_{r=1}^{m'} \|\mathbf{z}^{(r)} - \bar{z}^{(r)}\mathbf{1}\|_2^2 \leq \frac{1}{\lambda_2(\mathcal{G})} \sum_{r=1}^{m'} \mathbf{z}^{(r)\top} \mathbf{L} \mathbf{z}^{(r)} = \frac{E^{\text{dir}}(\mathbf{Z})}{\lambda_2(\mathcal{G})}, \quad (17)$$

which is (4). The constant $1/\lambda_2$ is sharp, attained by the Fiedler vector ψ_2 . The corresponding bound for the normalized Dirichlet energy $E_{\text{sym}}^{\text{dir}}(\mathbf{Z}) = \text{tr}(\mathbf{Z}^\top \Delta \mathbf{Z})$ is obtained by the change of variable $\mathbf{y}_v = \mathbf{Z}_v / \sqrt{d_v}$ and applying the same argument with $\lambda_2(\mathbf{L})$ in the bound. $\square$

B.3 Proof of Lemma 2 (Bi-Lipschitz Lower Bound for a GIN)

The block under analysis is one GIN aggregation followed by a two-layer MLP combine,

$$h_\theta(\mathbf{X}) = \sigma_2(\sigma_1(\mathbf{M}\mathbf{X}\mathbf{W}_1) \mathbf{W}_2), \quad (18)$$

acting on $\mathbf{X} \in \mathbb{R}^{N \times m}$ with $\mathbf{M} \in \mathbb{R}^{N \times N}$ mixing rows, and $\mathbf{W}_1 \in \mathbb{R}^{m \times m'}$, $\mathbf{W}_2 \in \mathbb{R}^{m' \times m''}$ acting on the feature (column) dimension. For leaky-ReLU activations (or ReLU with a fixed margin $\delta > 0$ along $\gamma(t) = \mathbf{X} + t\Delta\mathbf{X}$, ensuring no piecewise-linear region change), each σ_k has derivative $\geq \alpha$ everywhere along the path.

Vectorized Jacobian. Using $\text{vec}(\mathbf{A}\mathbf{Y}\mathbf{B}) = (\mathbf{B}^\top \otimes \mathbf{A})\text{vec}(\mathbf{Y})$,

$$\begin{aligned} \text{vec}(\mathbf{M}\mathbf{X}\mathbf{W}_1) &= (\mathbf{W}_1^\top \otimes \mathbf{M}) \text{vec}(\mathbf{X}), \\ \text{vec}(\sigma_1(\mathbf{M}\mathbf{X}\mathbf{W}_1) \mathbf{W}_2) &= (\mathbf{W}_2^\top \otimes \mathbf{I}_N) \mathbf{D}_1 (\mathbf{W}_1^\top \otimes \mathbf{M}) \text{vec}(\mathbf{X}), \end{aligned}$$

so the full Jacobian in vectorized form is

$$\mathbf{J}(\mathbf{X}) = \mathbf{D}_2 (\mathbf{W}_2^\top \otimes \mathbf{I}_N) \mathbf{D}_1 (\mathbf{W}_1^\top \otimes \mathbf{M}), \quad (19)$$

with $\mathbf{D}_k = \text{diag}(\sigma'_k) \succeq \alpha \mathbf{I}$.

575 **Singular-value chain.** For any product $\mathbf{ABCD}$ where each factor has $\sigma_{\min} > 0$, $\sigma_{\min}(\mathbf{ABCD}) \geq$
576 $\sigma_{\min}(\mathbf{A})\sigma_{\min}(\mathbf{B})\sigma_{\min}(\mathbf{C})\sigma_{\min}(\mathbf{D})$ (chain rule for singular values; equivalent to applying the lower
577 bound $\|\mathbf{Ax}\| \geq \sigma_{\min}(\mathbf{A})\|\mathbf{x}\|$ inductively). Using $\sigma_{\min}(\mathbf{P} \otimes \mathbf{Q}) = \sigma_{\min}(\mathbf{P})\sigma_{\min}(\mathbf{Q})$,

$$\begin{aligned}\sigma_{\min}(\mathbf{J}) &\geq \sigma_{\min}(\mathbf{D}_2) \sigma_{\min}(\mathbf{W}_2^\top \otimes \mathbf{I}_N) \sigma_{\min}(\mathbf{D}_1) \sigma_{\min}(\mathbf{W}_1^\top \otimes \mathbf{M}) \\ &\geq \alpha \cdot m_2 \cdot \alpha \cdot m_1 \kappa = \alpha^2 m_1 m_2 \kappa = c_{\text{lower}}.\end{aligned}$$

578 **Fundamental theorem of calculus.** For $\gamma(t) = \mathbf{X} + t \Delta \mathbf{X}$, the strict-margin assumption ensures
579 that no entry of the pre-activation $\mathbf{W}_1 \mathbf{M} \gamma(t)$ or $\mathbf{W}_2 \sigma_1(\mathbf{W}_1 \mathbf{M} \gamma(t))$ changes sign along $[0, 1]$, so $\mathbf{D}_1$
580 and $\mathbf{D}_2$ are *constant* in t . Hence $\mathbf{J}(\gamma(t)) \equiv \mathbf{J}$ is constant, and $h_\theta \circ \gamma$ is exactly affine on $[0, 1]$:

$$\text{vec}(h_\theta(\mathbf{X} + \Delta \mathbf{X}) - h_\theta(\mathbf{X})) = \mathbf{J} \text{vec}(\Delta \mathbf{X}). \quad (20)$$

581 Taking the ℓ_2 norm (= Frobenius norm in matrix form) and applying $\|\mathbf{J} \mathbf{x}\|_2 \geq \sigma_{\min}(\mathbf{J}) \|\mathbf{x}\|_2$,

$$\|h_\theta(\mathbf{X} + \Delta \mathbf{X}) - h_\theta(\mathbf{X})\|_F \geq c_{\text{lower}} \|\Delta \mathbf{X}\|_F. \quad (21)$$

582 For a true two-layer GIN $\sigma_2(\mathbf{M} \sigma_1(\mathbf{M} \mathbf{X} \mathbf{W}_1) \mathbf{W}_2)$, the second Kronecker factor in the chain becomes
583 $(\mathbf{W}_2^\top \otimes \mathbf{M})$ instead of $(\mathbf{W}_2^\top \otimes \mathbf{I}_N)$, contributing an extra $\sigma_{\min}(\mathbf{M}) \geq \kappa$ and giving $c_{\text{lower}} =$
584 $\alpha^2 m_1 m_2 \kappa^2$. $\square$

585 **Remark 1** (Beyond strict margin). Without the strict-margin condition γ may cross piecewise-linear
586 boundaries of σ_k . Then $h_\theta(\mathbf{X} + \Delta \mathbf{X}) - h_\theta(\mathbf{X}) = (\sum_k p_k \mathbf{J}_k) \Delta \mathbf{X}$ with $\mathbf{J}_k$ the per-region Jacobians
587 (each satisfying $\sigma_{\min}(\mathbf{J}_k) \geq c_{\text{lower}}$) and $p_k \geq 0, \sum p_k = 1$. Convex combinations of matrices
588 with $\sigma_{\min} \geq c$ do not in general satisfy the same lower bound, so c_{lower} may degrade across region
589 boundaries. We therefore use Lemma 2 only as a local non-degeneracy regularity statement in
590 Theorem 1; this is sufficient for our purposes.

591 B.4 Proof of Theorem 1 (Noise Forces Energy Growth)

592 Let $\mathbf{h}_i, \mathbf{h}_j \in \mathbb{R}^{N \times m'}$ be the post-message-passing per-node signals of $\mathcal{G}^i$ and $\mathcal{G}^j$ (we assume $N^i =$
593 $N^j = N$; the unequal- N case is treated in the Remark in the main text). Let $\mathbf{g}_i = \frac{1}{N} \sum_v \mathbf{h}_{i,v} \in \mathbb{R}^{m'}$
594 be the mean-pooled graph descriptor, and let $\bar{\mathbf{h}}_i = \mathbf{1}_N \mathbf{g}_i^\top$ be its replication to a constant per-node
595 signal. By the Frobenius-norm triangle inequality,

$$\|\mathbf{h}_i - \mathbf{h}_j\|_F \leq \|\mathbf{h}_i - \bar{\mathbf{h}}_i\|_F + \|\bar{\mathbf{h}}_i - \bar{\mathbf{h}}_j\|_F + \|\bar{\mathbf{h}}_j - \mathbf{h}_j\|_F. \quad (22)$$

596 **Middle term.** Because both $\bar{\mathbf{h}}_i, \bar{\mathbf{h}}_j$ are constant signals of size $N \times m'$,

$$\|\bar{\mathbf{h}}_i - \bar{\mathbf{h}}_j\|_F^2 = \sum_{v=1}^N \|\mathbf{g}_i - \mathbf{g}_j\|_2^2 = N \|\mathbf{g}_i - \mathbf{g}_j\|_2^2, \quad (23)$$

597 so $\|\bar{\mathbf{h}}_i - \bar{\mathbf{h}}_j\|_F = \sqrt{N} \|\mathbf{g}_i - \mathbf{g}_j\|_2$.

598 **Variance terms.** Identifying $\mathbf{h}_i$ with $\mathbf{Z}^i$ (the layer- L representation whose Dirichlet energy we
599 measure), Lemma 1 applied to $\mathbf{Z}^i$ gives

$$\|\mathbf{h}_i - \bar{\mathbf{h}}_i\|_F^2 = \sum_v \|\mathbf{h}_{i,v} - \mathbf{g}_i\|_2^2 \leq \frac{E^{\text{dir}}(\mathbf{Z}^i)}{\lambda_2(\mathcal{G}^i)}, \quad (24)$$

600 and analogously for $\mathcal{G}^j$. Taking square roots and combining,

$$\|\mathbf{h}_i - \mathbf{h}_j\|_F \leq \sqrt{N} \|\mathbf{g}_i - \mathbf{g}_j\|_2 + \sqrt{\frac{E^{\text{dir}}(\mathbf{Z}^i)}{\lambda_2(\mathcal{G}^i)}} + \sqrt{\frac{E^{\text{dir}}(\mathbf{Z}^j)}{\lambda_2(\mathcal{G}^j)}}, \quad (25)$$

601 which is (6).

602 **Lower bound from the noisy training objective.** The graph-level prediction is $\hat{\mathbf{y}} = R(\mathbf{h})$ for some
603 permutation-invariant readout R that is L -Lipschitz in the Frobenius norm (mean- or sum-pooling
604 followed by an MLP with bounded weights). At convergence, the noisy training objective enforces
605 $\hat{\mathbf{y}}_i$ to lie near a one-hot vector at c_i and $\hat{\mathbf{y}}_j$ near a one-hot at $c_j \neq c_i$, so $\|\hat{\mathbf{y}}_i - \hat{\mathbf{y}}_j\|_2 \geq \Delta_{\text{out}}$ for some
606 $\Delta_{\text{out}} > 0$. By the readout's upper-Lipschitz property,

$$\|\mathbf{h}_i - \mathbf{h}_j\|_F \geq L^{-1} \|\hat{\mathbf{y}}_i - \hat{\mathbf{y}}_j\|_2 \geq L^{-1} \Delta_{\text{out}} =: \Theta_{\text{noise}}. \quad (26)$$

The role of Lemma 2 is to make this lower bound non-vacuous: $c_{\text{lower}} > 0$ ensures that the encoder has a positive lower-Lipschitz constant at convergence, so the noisy-label constraint cannot be satisfied by collapsing $(\mathbf{X}^i, \mathbf{A}^i)$ and $(\mathbf{X}^j, \mathbf{A}^j)$ to the same internal representation; equivalently, the encoder remains in a regularity regime where $\|\mathbf{h}_i - \mathbf{h}_j\|_F$ is a meaningful quantity.

Combining upper and lower bounds. Under the structural-similarity assumption $\|\mathbf{g}_i - \mathbf{g}_j\|_2 \leq \varepsilon$, substituting $\Theta_{\text{noise}} \leq \|\mathbf{h}_i - \mathbf{h}_j\|_F$ into (6) and rearranging,

$$\sqrt{\frac{E^{\text{dir}}(\mathbf{Z}^i)}{\lambda_2(\mathcal{G}^i)}} + \sqrt{\frac{E^{\text{dir}}(\mathbf{Z}^j)}{\lambda_2(\mathcal{G}^j)}} \geq \Theta_{\text{noise}} - \sqrt{N} \varepsilon, \quad (27)$$

which is (7). $\square$

Unequal graph sizes. When $N^i \neq N^j$, set $N = \max(N^i, N^j)$ and pad the smaller graph by appending $N - N^{\min}$ rows equal to its mean $\mathbf{g}_\bullet$, with the corresponding padded vertices isolated (no edges). Two facts make this padding harmless: (i) appending $\mathbf{g}_\bullet$ -valued rows leaves the graph mean unchanged, and (ii) isolated padded vertices contribute nothing to the Dirichlet energy, so $E^{\text{dir}}(\mathbf{Z}^{\text{pad}}) = E^{\text{dir}}(\mathbf{Z})$. The within-graph variance $\|\mathbf{h}_\bullet^{\text{pad}} - \bar{\mathbf{h}}_\bullet^{\text{pad}}\|_F^2 = \sum_{v=1}^{N^{\min}} \|\mathbf{h}_{\bullet,v} - \mathbf{g}_\bullet\|^2$ is also preserved, so Lemma 1 applies unchanged. The middle term becomes $\|\bar{\mathbf{h}}_i^{\text{pad}} - \bar{\mathbf{h}}_j^{\text{pad}}\|_F = \sqrt{N} \|\mathbf{g}_i - \mathbf{g}_j\|_2$, and the readout’s lower bound is preserved up to a fixed pooling constant. Hence (6) and (7) hold with $N \rightarrow \max(N^i, N^j)$.

B.5 Proof of Lemma 3 (Spectral Counting)

The combinatorial Laplacian $\mathbf{L}^i$ is symmetric positive semi-definite. By Anderson–Morley’s bound, every eigenvalue λ_j^i of $\mathbf{L}^i$ satisfies $\lambda_j^i \leq \max_{(u,v) \in \mathcal{E}^i} (d_u + d_v) \leq 2d_{\max}(\mathcal{G}^i)$. Equivalently, the normalized Laplacian $\Delta^i = \mathbf{D}^{-1/2} \mathbf{L}^i \mathbf{D}^{-1/2}$ has eigenvalues in $[0, 2]$, with $\lambda_{\max}(\Delta^i) = 2$ iff $\mathcal{G}^i$ has a bipartite component [Chung, 1997, Lemma 1.7, Thm. 1.13]. The number of distinct eigenvalues of $\mathbf{L}^i$ is at most N^i .

For any threshold $\tau > 0$, define $n^i(\tau) = |\{j : \lambda_j^i \geq \tau\}|$. We claim $n^i(\tau)$ is non-decreasing under (a) edge addition and (b) vertex addition with at least one incident edge.

(a) *Edge addition:* if $\mathcal{G}^{i'} = \mathcal{G}^i + \{(u, v)\}$, then $\mathbf{L}^{i'} = \mathbf{L}^i + (\mathbf{e}_u - \mathbf{e}_v)(\mathbf{e}_u - \mathbf{e}_v)^\top$, a rank-one PSD update. By Weyl’s monotonicity theorem (Horn and Johnson, 2012, Cor. 4.3.12), $\lambda_k(\mathbf{L}^{i'}) \geq \lambda_k(\mathbf{L}^i)$ for every k . Hence if $\lambda_j^i \geq \tau$ then $\lambda_j^{i'} \geq \tau$, so $n^{i'}(\tau) \geq n^i(\tau)$.

(b) *Vertex addition:* write $\mathcal{G}^{i'}$ as the result of starting from the padded graph $\mathcal{G}^i \sqcup \{v_{N+1}\}$ (an isolated extra vertex) and then adding all incident edges $\{(v_{N+1}, u) : u \in \mathcal{N}(v_{N+1})\}$ one at a time. The padded Laplacian is $\mathbf{L}_0 = \mathbf{L}^i \oplus [0] \in \mathbb{R}^{(N^i+1) \times (N^i+1)}$, with spectrum $\Lambda(\mathbf{L}^i) \cup \{0\}$. Each subsequent edge addition is a rank-one PSD update covered by case (a), so $\lambda_k(\mathbf{L}^{i'}) \geq \lambda_k(\mathbf{L}_0)$ for all k . Since $\Lambda(\mathbf{L}_0) \setminus \{0\} = \Lambda(\mathbf{L}^i) \setminus \{0\}$, the count above any $\tau > 0$ satisfies $n^{i'}(\tau) \geq n_0(\tau) = n^i(\tau)$. (The case $\tau = 0$ is trivial since $n^{i'}(0) = N^i + 1 \geq N^i = n^i(0)$.)

Now let $\bar{\mathbf{L}} = \text{blkdiag}(\mathbf{L}^1, \dots, \mathbf{L}^n)$ as in Proposition 1. By the block-diagonal structure, every eigenpair $(\bar{\lambda}_u, \bar{\psi}_u)$ of $\bar{\mathbf{L}}$ is supported on exactly one $\mathcal{G}^i$. Sharpening that lifts spectral mass onto eigenvectors with $\bar{\lambda}_u \geq \tau$ therefore allocates a share to $\mathcal{G}^i$ that is bounded above by $n^i(\tau) / \sum_k n^k(\tau)$, which grows with N^i and connection density of $\mathcal{G}^i$. $\square$

B.6 Proof of Corollary 1 (Asymmetric Burden)

Suppose $H(\mathcal{G}^j) \geq H(\mathcal{G}^i)$ either via more nodes at fixed sparsity (Theorem 1 of Qian et al., 2025) or via denser connectivity at fixed $|V|$ (Theorem 2 of Qian et al., 2025). Lemma 3 then implies $n^j(\tau) \geq n^i(\tau)$ for every τ , where $n^i(\tau)$ counts eigenvalues of $\mathcal{G}^i$ above τ . Now consider any sharpening update of $\bar{\mathbf{Z}}$ that increases the spectral coefficients on eigenpairs with $\bar{\lambda}_u \geq \tau$. By Proposition 1, each such eigenpair lies entirely on one of $\mathcal{G}^i$ or $\mathcal{G}^j$, and the count of available high-frequency eigenvectors on $\mathcal{G}^j$ is at least that on $\mathcal{G}^i$. Hence on average a $\bar{\lambda}_u \geq \tau$ component is more likely to be supported on $\mathcal{G}^j$, and the resulting energy increase is preferentially absorbed by $\mathcal{G}^j$. By Section 3.3 of Qian et al. [2025], increasing $E^{\text{dir}}(\mathbf{Z}^j)$ along its high-frequency eigenvectors collapses the pairwise distinctness of node-feature distributions inside $\mathcal{G}^j$, the precise quantity that pooled

graph classifiers rely on for separation. Hence high-entropy graphs lose discriminability faster than low-entropy graphs. $\square$

B.7 Why a Decrease of $E_{\text{set}}^{\text{dir}}$ Is Non-Trivial in Graph Classification

Three consecutive observations together justify the term “non-trivial”:

(i) Noise is injected at the graph-label level, not at the node-feature level. Therefore, the increase of $E^{\text{dir}}(\mathbf{Z}^i)$ during noise fitting cannot be a re-encoding of the noise itself; it must be a behavioral response of the network through training dynamics.

(ii) For a randomly-initialized network, $E_{\text{set}}^{\text{dir}}$ is independent of class labels by construction. Any correlation between $E_{\text{set}}^{\text{dir}}$ and labels emerges from the optimization process, not from the data.

(iii) By Theorem 1 and Corollary 1, the only optimization path that fits incorrect graph-level labels is to lift spectral mass into high-entropy directions of $\bar{\mathbf{L}}$, which are precisely the directions that a clean, generalization-driven network should preserve. Constraining $E_{\text{set}}^{\text{dir}}$ therefore selectively blocks the noise-fitting path while allowing the clean-learning path, an alignment that does *not* hold for node classification, where label noise corrupts the same nodes whose smoothness the energy measures, and where decreasing energy can erase the discriminative signal.

These three points collectively make the connection between energy reduction and noise robustness in graph classification a non-trivial property. The analogous statement in node classification reduces to “the model fits homophily, and homophily is what we measured.”

C Experimental Setting

C.1 Datasets

ogbg-ppa [Szkarczyk et al., 2018a]: 158,100 protein-association graphs across 37 taxonomic groups, derived from STRING [Szkarczyk et al., 2018b]; the standard OGB graph-property-prediction benchmark. Modified PPA-30%-6 takes 30% of training graphs from 6 selected classes, used as a controlled testbed. **PROTEINS** [Borgwardt et al., 2005]: 1,113 protein graphs; 2 classes; mean 39 nodes/graph. **ENZYMES** [Borgwardt et al., 2005]: 600 enzyme graphs across 6 EC top-level classes. **MSRC-21** [Neumann et al., 2016]: 563 image-segmentation graphs, 20 classes. **MUTAG** [Kriege and Mutzel, 2012]: 188 molecular graphs, 2 classes. **IMDB-BINARY** [Yanardag and Vishwanathan, 2015b]: 1,000 social graphs, 2 classes. **REDDIT-MULTI-12K** [Yanardag and Vishwanathan, 2015b]: 11,929 thread-graphs, 11 classes. **MNIST-graph**: 55,000 superpixel graphs over the MNIST images, 10 classes.

C.2 Synthetic Datasets for Failure-Mode F2

We sample average graph orders $\bar{N} \in \{5, 10, \dots, 60\}$, drawing actual node counts from $\text{Poisson}(\bar{N})$. Each dataset has 6 classes with 1,400 graphs/class, fixed average degree 2, edges sampled uniformly. Node features are $\mathcal{N}(\mu_c, 1.5)$ with $\mu_c = f(\text{class})$. Symmetric noise rate is fixed at 35%.

C.3 Architectures and Hyperparameters

Unless otherwise stated, we use a 5-layer GIN [Xu et al., 2019] or GCN [Kipf and Welling, 2017] with 300 hidden units, Adam ($\eta = 10^{-3}$, weight decay 10^{-4}), batch size 32, 200–1000 epochs, $\eta_u = 1$ for GCOD’s outlier parameter. Compute: NVIDIA RTX A6000.

C.4 Spectral Bias Implementation Details (CE+W2)

For each layer l , after a gradient update, decompose $\mathbf{W}_l^2 = \Phi_l^2 \mu_l^2 (\Phi_l^2)^{-1}$ and replace by $\Phi_l^2 [\mu_l^2]^+ (\Phi_l^2)^{-1}$, where $[\cdot]^+$ is element-wise ReLU on the eigenvalues. The projection is detached from the gradient graph. We apply it to the post-aggregation matrix $\mathbf{W}^2$ of every layer; applying it also to $\mathbf{W}^1$ marginally improves smoothing but degrades stability. Convergence variance under this scheme is shown in Fig. 8.

697 C.5 Class-Specific and Fixed Bounds for $\mathcal{L}_{DE}$

698 **Class-specific:** after every epoch, set $U_c = |\mathcal{D}_c^{\text{val}}|^{-1} \sum_{i \in \mathcal{D}_c^{\text{val}}} E^{\text{dir}}(\mathbf{Z}^i)$. The validation set must be
 699 clean to keep U_c class-specific. **Fixed:** $U_c = U$ globally. We progressively decrease U from a large
 700 initial value until accuracy under clean labels begins to drop; the final U is the largest value that
 701 achieves a noticeable energy ceiling without over-smoothing.

702 C.6 Design of GCOD

703 We adopt the centroid-based soft labels of Wani et al. [2023]: normalize $\phi(\mathbf{z}_i) = \mathbf{z}_i / \|\mathbf{z}_i\|$; per-
 704 class centroid $\phi_c = \bar{\mathbf{z}}_c / \|\bar{\mathbf{z}}_c\|$; soft target $\tilde{\mathbf{y}}_i = [\max(\phi(\mathbf{z}_i)\phi_c^\top, 0)\mathbf{y}_i]$. Equations (9)–(11) are then
 705 optimized as in the main text. The training-accuracy feedback in $\mathcal{L}_1$ and $\mathcal{L}_3$ is critical: without it, $\mathbf{u}$
 706 dominates early in training and discounts samples before genuine patterns are learned (Fig. 7).

707 D Failure-Mode Experiments (Extended)

708 Additional plots for the F1–F3 failure modes are provided here for completeness, including effect-of-
 709 dataset-size for both PPA and synthetic datasets.




Figure 3: Effect of training-set size on noise memorization (training accuracy on noisy labels only).
 (a) PPA subsampling at 20% symmetric noise. (b) Synthetic graphs of order 7, 6 classes, 35% noise.

710 E Dirichlet-Energy Amplification in High-Frequency Components

711 We use the HLFF-GNN framework [Xu et al., 2024] (implemented as FGRLConv) to decompose
 712 representations into low-frequency Z_1 , high-frequency Z_2 , and residual Y . Under the composite loss


$$\mathcal{L} = \|Y - X\|_F^2 + \lambda \text{tr}(Z_1^\top \Delta Z_1) + \alpha \text{tr}(Z_2^\top (\mathbf{I} - \Delta) Z_2) + \beta (\|Y^\top Z_1\|_F^2 + \|Y^\top Z_2\|_F^2), \quad (28)$$

713 the spectral form of the high-frequency energy is $E^{\text{dir}}(Z_2) = \sum_{r,u} \lambda_u (\psi_u^\top Z_{2,r})^2$, in which large λ_u
 714 amplify any noise-driven content. On ENZYMES with 30% symmetric noise, we observe a sharp
 715 monotone increase in $E^{\text{dir}}(Z_2)$, while $E^{\text{dir}}(Z_1)$ and $E^{\text{dir}}(Y)$ remain stable, consistent with Theorem 1
 716 and Corollary 1.


717 F Additional Empirical Studies on Energy Behavior

718 **Energy under clean overfitting vs. noise.** On ENZYMES without regularization, a deliberately
 719 overfit GIN exhibits rising training accuracy and rising $E_{\text{set}}^{\text{dir}}$, confirming that energy tracks overfitting
 720 in general. However, the rise under noise is markedly steeper.

721 **Energy scales with noise rate.** On the controlled PPA subset, we sweep symmetric noise $\in$
 722 $\{10, 20, 30, 40, 60\}\%$. All curves follow the same U-shape (initial decrease, then rise), and the
 723 asymptotic energy increases monotonically with noise rate. This monotonicity is what makes $E_{\text{set}}^{\text{dir}}$ a
 724 diagnostic.




(a) Normalized $E^{\text{dir}}(Y, Z_1, Z_2)$.




(b) $E^{\text{dir}}(Z_2)$ alone.

Figure 4: HLFF-GNN frequency decomposition on ENZYMES. Solid lines: clean labels. Dashed: 30% symmetric noise. The high-frequency component Z_2 is the only one whose energy diverges under noise, supporting the theoretical prediction.



(a) Clean overfitting on ENZYMES.



(b) $E_{\text{set}}^{\text{dir}}$ vs. noise rate on PPA.

Figure 5: (a) Even without label noise, energy rises during overfitting; the surge under noise is much steeper. (b) Higher noise rates yield monotonically higher final energy.

725 G Additional Tables and Figures

Table 5: Sensitivity of GCOD to the learning rate η_u of the outlier parameter $\mathbf{u}$, under 40% asymmetric noise.

Dataset	$\eta_u = 1$	$\eta_u = 0.1$	% Change
PROTEINS	76.19	75.89	−0.39%
MNIST	72.61	72.48	−0.18%
ENZYMES	69.81	68.33	−2.12%
IMDB-B	72.89	71.94	−1.30%
MUTAG	93.19	92.61	−0.62%
REDDIT	47.94	48.08	+0.29%
MSRC-21	95.57	95.45	−0.13%

726 H Full ogbg-ppa: Do-No-Harm Check

727 The standard ogbg-ppa benchmark has 37 classes and $\sim 158,000$ graphs of mean order ≈ 240 , so a
 728 5-layer GIN with 300 hidden units is structurally *under-parameterized* for the task. We trained CE,
 729 $\mathcal{L}_{DE}$ (Fixed), CE+W2 and GCOD on the full benchmark with 0%, 20% and 40% symmetric noise.
 730 Two qualitative observations are robust across runs:

- 731 1. *Training accuracy on the noisy-labeled subset stays close to chance for all methods, including*
 732 *standard CE.* The under-parameterized model never reaches a regime where it can memorize the
 733 corrupted graph-level labels, consistent with our failure-mode analysis of Section 4: F1 (number
 734 of classes), F2 (graph order) and F3 (training-set size) all push the full PPA out of the noise-fitting
 735 regime.

Table 6: Performance of CE vs. GCOD with 20% symmetric label noise. Last column reports the absolute gain (GCOD−CE) on test accuracy.

Dataset	Metric	0% CE	20% CE	20% GCOD	20% (GCOD−CE)
MNIST	Best	72.69	66.30	71.26	+4.96
	Last	69.80	53.64	67.88	+14.24
	Diff	2.89	12.66	3.38	
ENZYMES	Best	73.33	64.16	68.69	+4.53
	Last	65.78	57.50	62.54	+5.04
	Diff	7.55	6.66	6.15	
MSRC-21	Best	96.69	90.26	94.69	+4.43
	Last	93.80	79.64	90.26	+10.62
	Diff	2.89	10.62	4.43	
PROTEINS	Best	81.16	76.23	79.38	+3.15
	Last	79.18	62.32	78.12	+15.80
	Diff	1.98	13.91	1.26	
MUTAG	Best	94.73	89.47	90.01	+0.54
	Last	84.21	68.42	86.84	+18.42
	Diff	10.52	21.05	3.17	
IMDB-B	Best	76.50	75.00	75.40	+0.40
	Last	71.60	71.00	73.50	+2.50
	Diff	4.90	4.00	1.90	
REDDIT	Best	48.15	45.05	44.98	−0.07
	Last	46.01	44.67	44.89	+0.22
	Diff	2.14	0.38	0.09	




Figure 6: (a) Train accuracy on known noisy and clean PPA subsets (4-class, 40% symmetric noise) for GCN with CE: the model never separates noisy from clean training samples without a robust loss. (b) Train/test accuracy on ENZYMES with GIN under clean and 20% symmetric noise.

736 2. *Test accuracy of our three methods is statistically indistinguishable from CE at every noise rate.*
737 None of CE+W2, $\mathcal{L}_{DE}$, or GCOD perturbs the underlying training dynamics enough to harm
738 clean generalization, since their corrective signals (negative-eigenvalue projection, energy upper
739 bound, outlier discounting) are only triggered by the high-frequency surge that does not occur on
740 the full PPA.

741 This is by design: the methods are conditional regularizers, not unconditional smoothers. The
742 take-away is that they impose *no measurable overhead in accuracy* when the underlying network is
743 already noise-robust, while providing substantial protection when over-parameterization, low graph
744 order, or small training sets push the network into the noise-fitting regime (Section 8, controlled
745 PPA-30%-6 subset). A fair generality claim therefore rests on the seven additional benchmarks in
746 Tables 2 and 6, with the full PPA serving exclusively as a do-no-harm sanity check.

747 I Cross-Setting Transfer: Cora Node Classification

748 To assess whether the energy perspective captures a property of GNN training under label noise that
749 is more general than the graph-classification setting it was designed for, we apply CE+W2 and GCOD




Figure 7: Ablation on the training-accuracy feedback in GCOD. Without scaling u by a_{train} (“GCOD WOT”), the outlier parameter dominates early and harms generalization; the full GCOD activates outlier-discounting gradually.




Figure 8: Test accuracy on PPA-30%-6 for CE and CE+W2. CE+W2 reaches comparable peak accuracy but exhibits unstable convergence due to the per-epoch eigendecomposition.

750 *unchanged* to node classification on Cora at 20% symmetric label noise. We use both GAT and GCN
 751 backbones, with the standard public split, and compare against 13 node-classification baselines that
 752 include CE, GCE [Ghosh et al., 2017], NRGNN [Dai et al., 2021], RTGNN [Yuan et al., 2023a],
 753 GNN-Cleaner, ERASE, CR-GNN, and the standard noise-robust losses adapted to node classification.

Table 7: Cross-setting transfer of GCOD and CE+W2 to node classification on Cora under 20% symmetric label noise. Numbers are test accuracy; ranks are with respect to the full 13-method comparison reported in the rebuttal materials. We report only the verified results for our two methods on each backbone; the “rank” field indicates their position in the full ranking. The methods are not modified, regularization parameters are kept at their graph-classification defaults.

Backbone	Method	Test Acc.	Rank (out of 13)
GAT	CE+W2	0.8157	#2
GAT	GCOD	0.7910	#4
GCN	GCOD	0.7403	#4
GCN	CE+W2	0.7247	#5

754 Both methods land in the top half of the ranking on each backbone, despite using neither homophily,
 755 contrastive losses, pseudo-labels, nor any node-specific machinery. CE+W2 in particular is competi-
 756 tive with the strongest specialized node-classification baselines on GAT. The interpretation, consistent
 757 with Theorem 1, is that the spectral mechanism by which noise fitting forces within-graph energy to
 758 grow is not specific to graph classification: in node classification, individual labels are also embedded
 759 in node-level representations whose Dirichlet energy reflects local smoothness, and constraining that
 760 smoothness penalizes the same noise-fitting path. The transfer is therefore not a coincidence but a
 761 corollary of the energy-based framing.




Figure 9: Average test accuracy and average within-graph $E_{\text{set}}^{\text{dir}}$ on MUTAG (0% noise) with GIN: classification accuracy improves while energy decays, the clean-data signature of a healthy GNN.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and Section 1 state precisely the claims supported by the theoretical results in Section 6 and the experiments in Section 8.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: A dedicated paragraph in the Conclusion discusses the absence of a closed-form noise-energy characterization, reliance on synthetic noise, and the lack of standardized real-world noisy graph-classification benchmarks.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: Proposition 1, Lemmas 1–2, Theorem 1, Lemma 3 and Corollary 1 state their assumptions explicitly in the main text, and full proofs are provided in Appendix B.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper?

Answer: [\[Yes\]](#)

Justification: Appendix C provides datasets, splits, hyperparameters, architectures, and the implementation details of all three methods. Anonymized code is provided in the supplemental material.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results?

Answer: [\[Yes\]](#)

Justification: All datasets are public benchmarks; an anonymous code repository accompanies the submission.

6. Experimental setting/details

Question: Does the paper specify all the training and test details necessary to understand the results?

Answer: [\[Yes\]](#)

Justification: Section 8 and Appendix C report all training and evaluation details.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [\[Yes\]](#)

Justification: All main results report mean $\pm$ std over 4 independent runs with different seeds.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources?

Answer: [\[Yes\]](#)

Justification: Appendix C reports the use of NVIDIA RTX A6000 GPUs and Table 4 reports relative training-time overheads.

810	9. Code of ethics
811	Question: Does the research conducted in the paper conform, in every respect, with the
812	NeurIPS Code of Ethics?
813	Answer: [Yes]
814	Justification: The work uses public benchmarks, does not involve human subjects, and the
815	broader-impact paragraph discusses application contexts.
816	10. Broader impacts
817	Question: Does the paper discuss both potential positive societal impacts and negative
818	societal impacts of the work performed?
819	Answer: [Yes]
820	Justification: Discussed in the Conclusion. The work primarily benefits scientific applica-
821	tions (chemistry, biology) where label noise is systematic.
822	11. Safeguards
823	Question: Does the paper describe safeguards that have been put in place for responsible
824	release of data or models that have a high risk for misuse?
825	Answer: [N/A]
826	Justification: The paper does not release new datasets or pretrained models with misuse risk.
827	12. Licenses for existing assets
828	Question: Are the creators or original owners of assets, properly credited and are the license
829	and terms of use explicitly mentioned and properly respected?
830	Answer: [Yes]
831	Justification: All datasets are cited; OGB, TU-Datasets and PyTorch-Geometric implemen-
832	tations are used under their respective licenses.
833	13. New assets
834	Question: Are new assets introduced in the paper well documented and is the documentation
835	provided alongside the assets?
836	Answer: [Yes]
837	Justification: The released code provides documentation for the proposed methods.
838	14. Crowdsourcing and research with human subjects
839	Question: For crowdsourcing experiments and research with human subjects, does the paper
840	include the full text of instructions given to participants and screenshots, if applicable, as
841	well as details about compensation (if any)?
842	Answer: [N/A]
843	Justification: No human-subject or crowdsourcing experiments.
844	15. Institutional review board (IRB) approvals or equivalent for research with human
845	subjects
846	Question: Does the paper describe potential risks incurred by study participants, whether
847	such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
848	approvals (or an equivalent approval/review based on the requirements of your country or
849	institution) were obtained?
850	Answer: [N/A]
851	Justification: No human subjects.
852	16. Declaration of LLM usage
853	Question: Does the paper describe the usage of LLMs if it is an important, original, or
854	non-standard component of the core methods in this research?
855	Answer: [N/A]
856	Justification: LLMs are not part of the core methods.